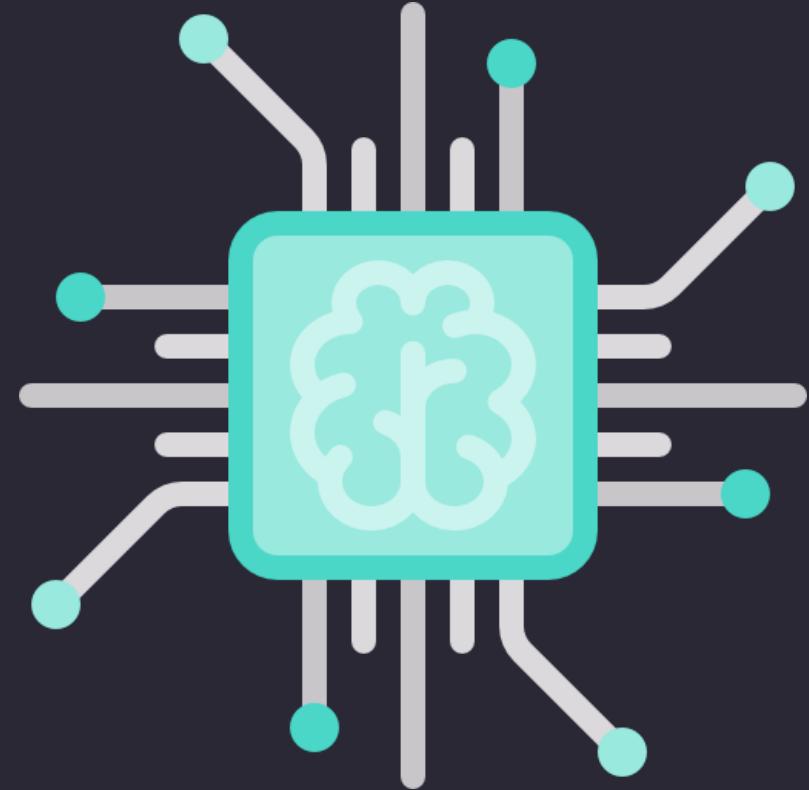


Deep Dissection

Understanding the ins and outs of Neural Networks

by Diogo Pinto



The road
ahead

The road ahead

Motivation

- Why deep learning
- Playground

The road ahead

Motivation

- Why deep learning
- Playground

Fundamentals

- Derivatives & Chain rule
- Perceptron
- Sigmoid function & Cross-entropy

The road ahead

Motivation

- Why deep learning
- Playground

Fundamentals

- Derivatives & Chain rule
- Perceptron
- Sigmoid function & Cross-entropy

The Internals

- Forward propagation
- Error computation
- Back propagation
- Parameters update

The road ahead

Motivation

- Why deep learning
- Playground

Fundamentals

- Derivatives & Chain rule
- Perceptron
- Sigmoid function & Cross-entropy

The Internals

- Forward propagation
- Error computation
- Back propagation
- Parameters update

Coffee Break



The road ahead

Motivation

- Why deep learning
- Playground

Fundamentals

- Derivatives & Chain rule
- Perceptron
- Sigmoid function & Cross-entropy

The Internals

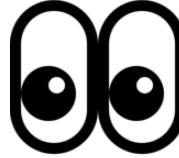
- Forward propagation
- Error computation
- Back propagation
- Parameters update

Coffee Break



Implementation

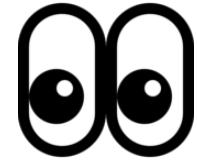
- Binary Classification – Census Income dataset

Who are you people? 

Who are you people?

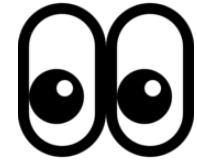
- Who is used to work with Python?

Who are you people?



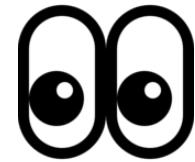
- Who is used to work with Python?
- Who already did something **close to this**?

Who are you people?



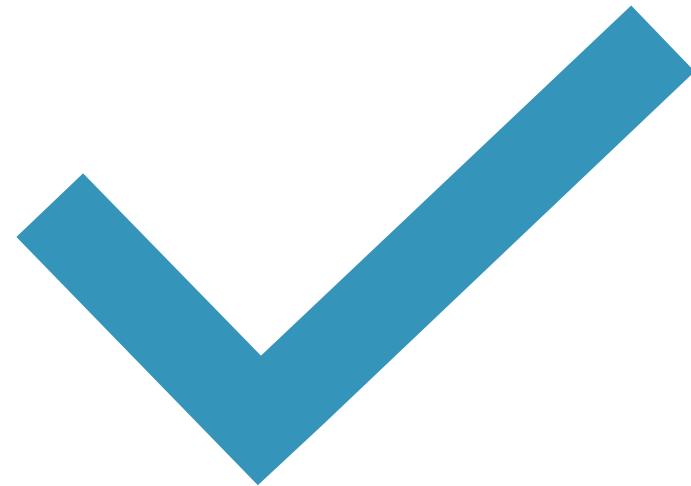
- Who is used to work with Python?
- Who already did something close to this?
- Who likes **matrices and letters** in the place of numbers?

Who are you people?



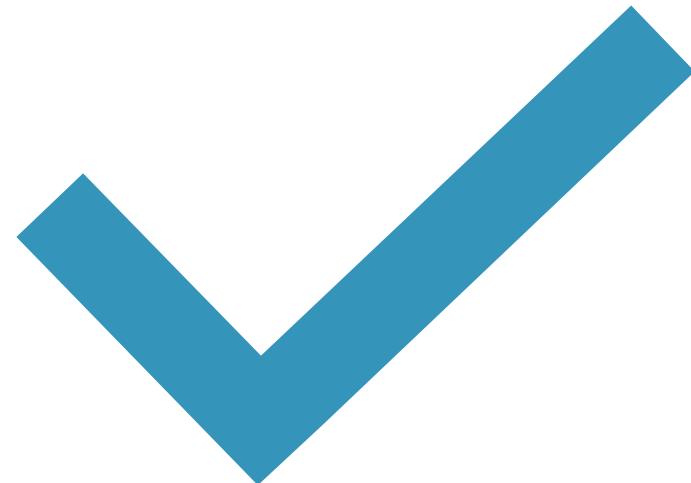
- Who is used to work with Python?
- Who already did something close to this?
- Who likes matrices and letters in the place of numbers?
- Who had a good night of **sleep**?

Assumptions



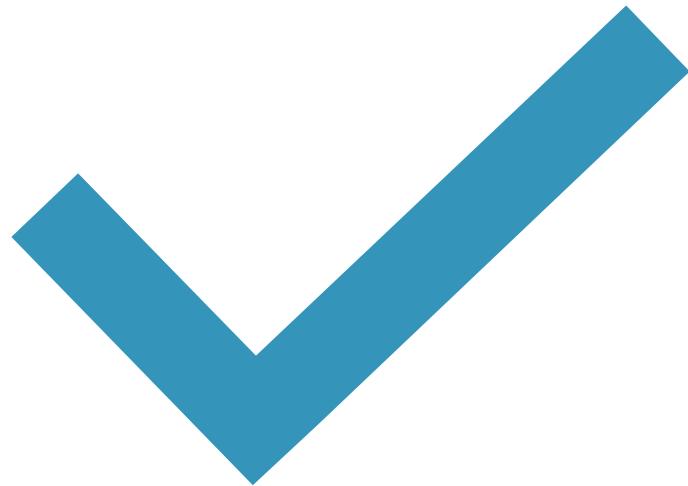
Assumptions

- A bit of **Algebra** and **Calculus**



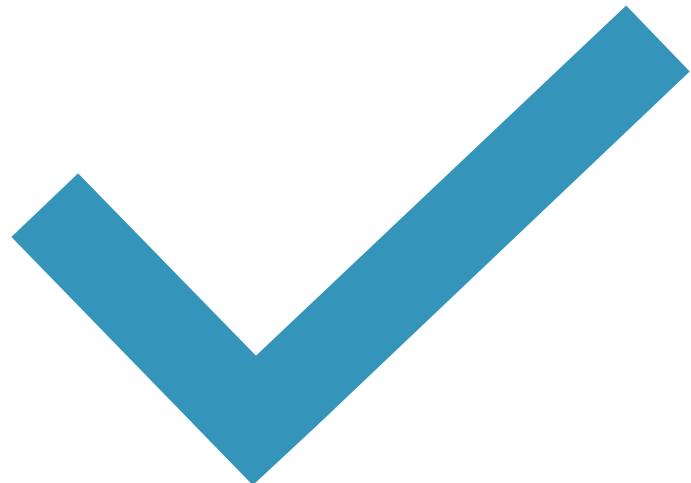
Assumptions

- A bit of **Algebra and Calculus**
- Introductory level of **Machine Learning** knowledge



Assumptions

- A bit of **Algebra** and **Calculus**
- Introductory level of **Machine Learning** knowledge



Difficulty can be calibrated, give feedback!

Motivation

Why am I here?

"Don't limit your challenges, challenge your limits"

Unknown

The trick is in the representation...

The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive

The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive

01100
10110
11110

The trick is in the representation...

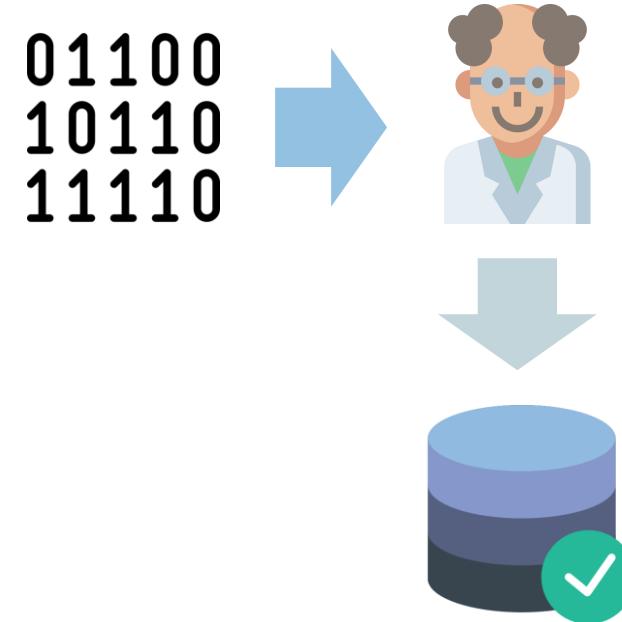
- Machine learning is traditionally **feature engineering** intensive

01100
10110
11110



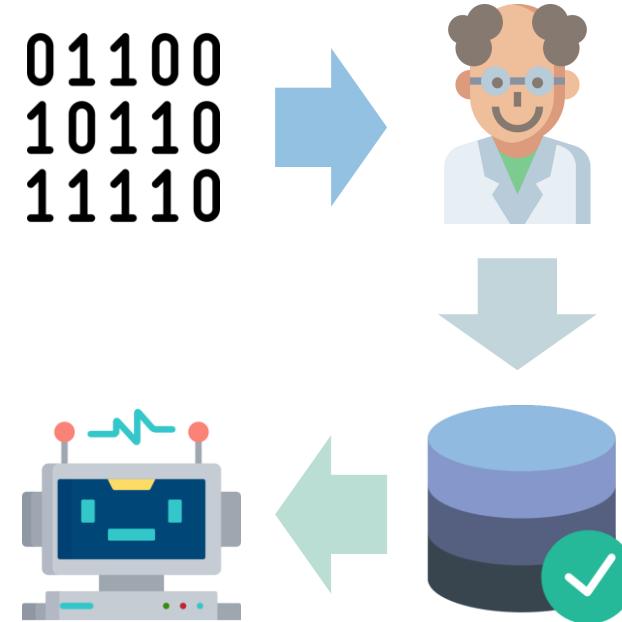
The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive



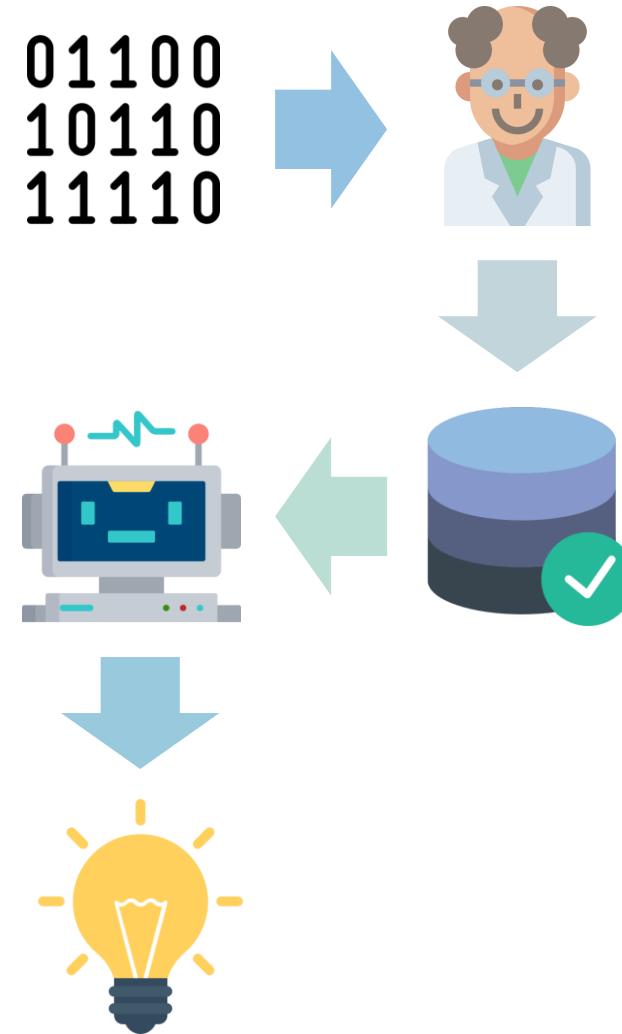
The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive



The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive

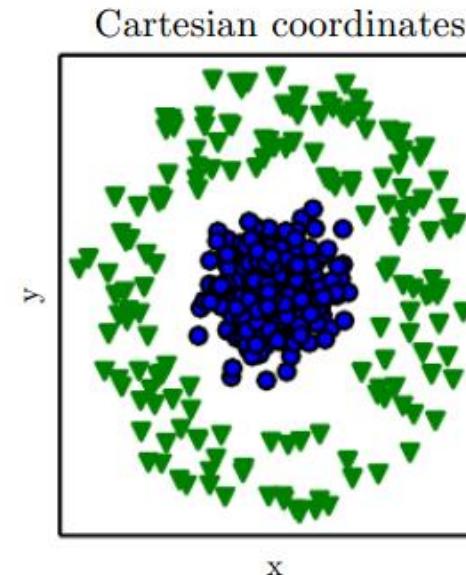


The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**

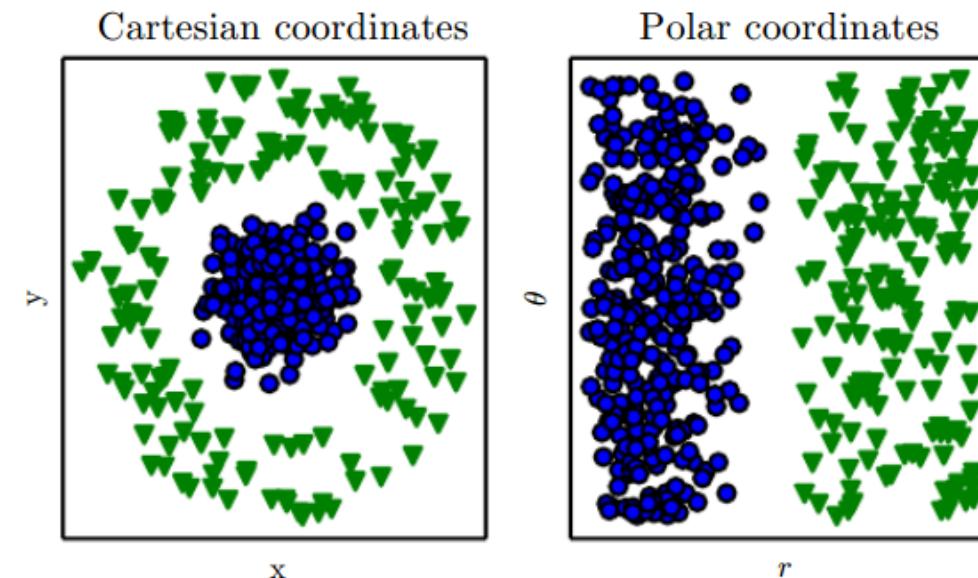
The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**



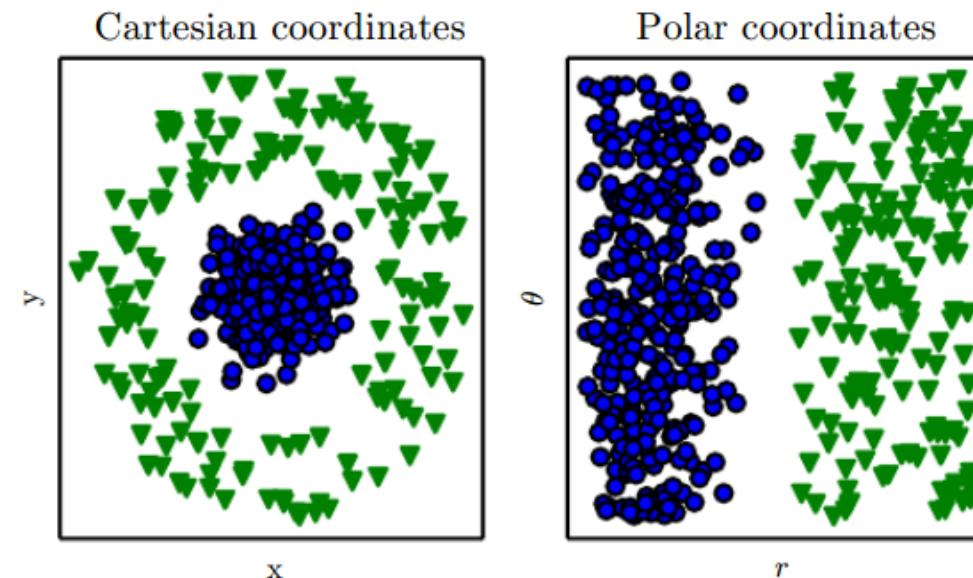
The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**



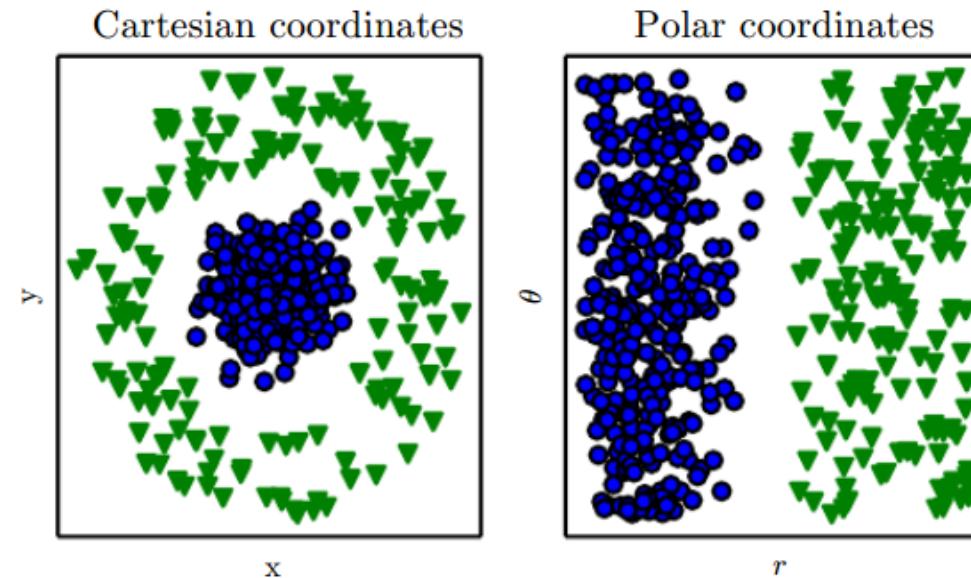
The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**
 - Disentangle them



The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**
 - Disentangle them
 - Discard op the unnecessary ones

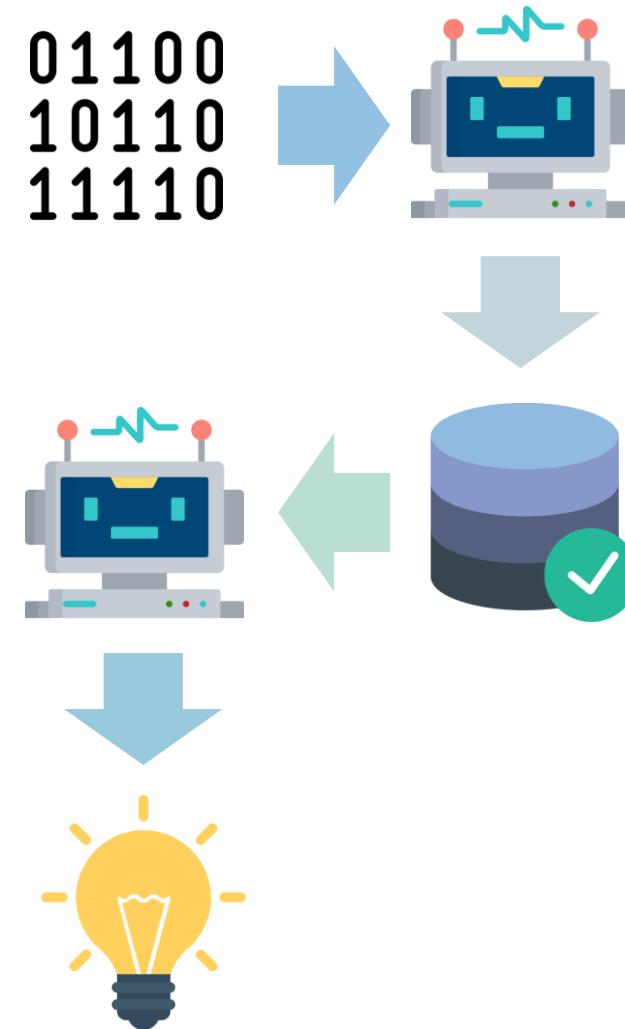


The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**
 - Disentangle them
 - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process

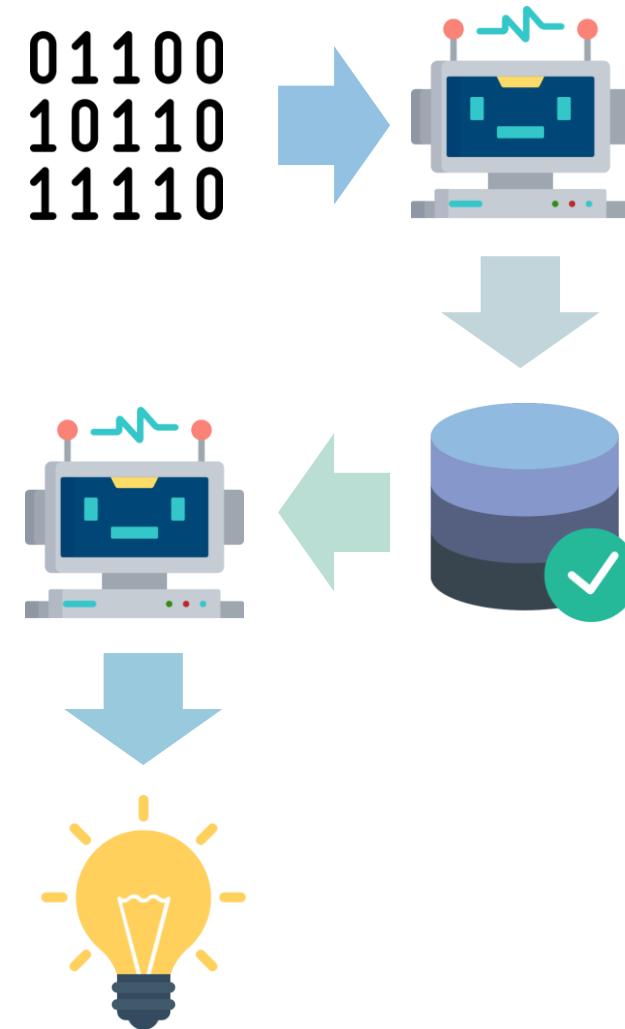
The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**
 - Disentangle them
 - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process



The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**
 - Disentangle them
 - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process
 - Learning **high-level concepts** can be as hard as the original problem

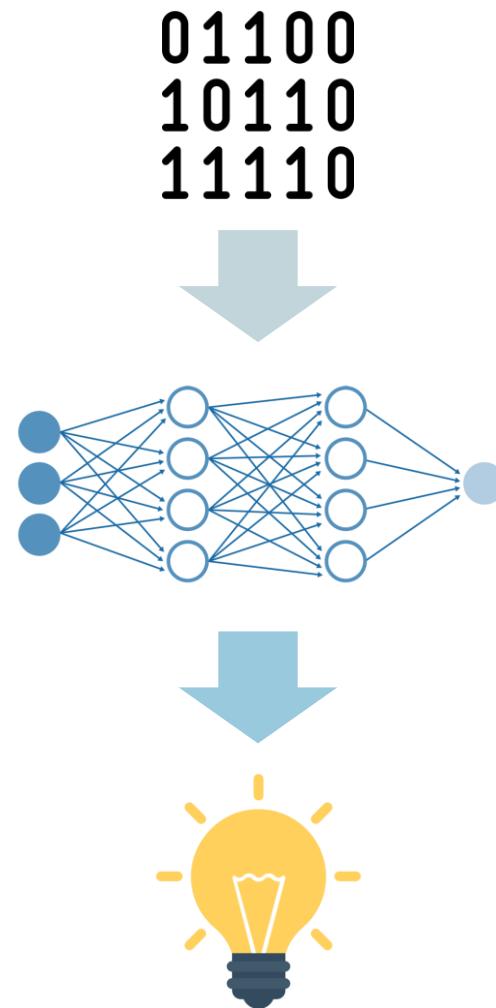


The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**
 - Disentangle them
 - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process
 - Learning **high-level concepts** can be as hard as the original problem
- **Deep learning** tackles this heads on

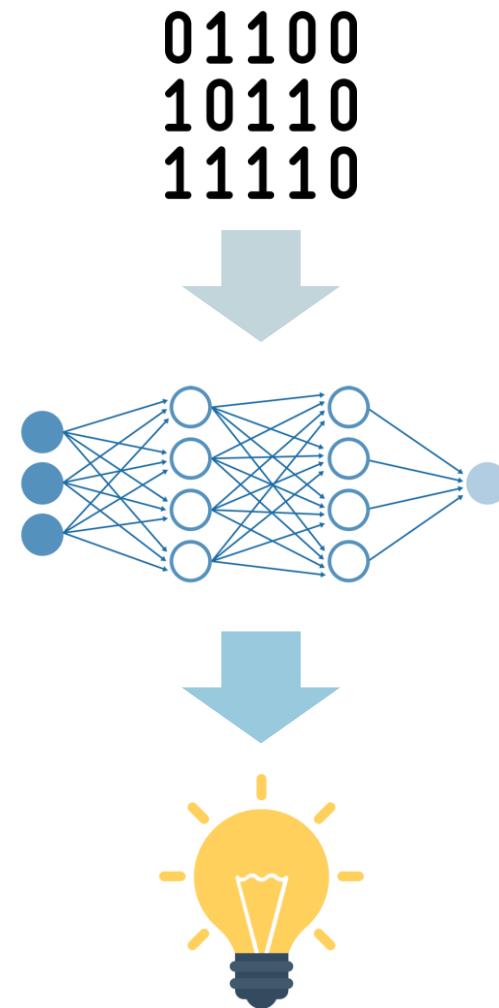
The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**
 - Disentangle them
 - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process
 - Learning **high-level concepts** can be as hard as the original problem
- **Deep learning** tackles this heads on



The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
 - Identifying the **sources of influence**
 - Disentangle them
 - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process
 - Learning **high-level concepts** can be as hard as the original problem
- **Deep learning** tackles this heads on
 - **Hierarchical combination** of simpler concepts into more complex ones



A quick visual demo...

[Neural Networks Playground](#)

Fundamentals

What do I need to know?

"You can't build a great building on a weak foundation."

Gordon B. Hinckley

It's Calculus Time



It's Calculus Time



- Derivatives

It's Calculus Time



- Derivatives
 - $y = f(x)$

It's Calculus Time

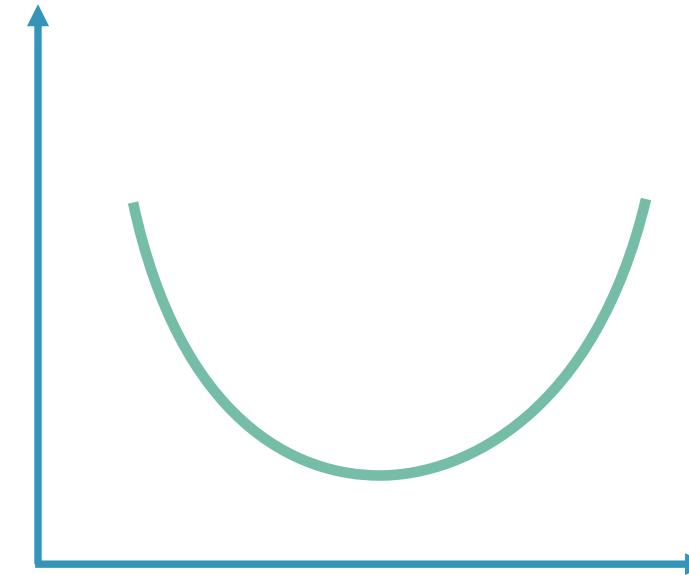


- **Derivatives**
 - $y = f(x)$
 - How a small change in x changes y

It's Calculus Time



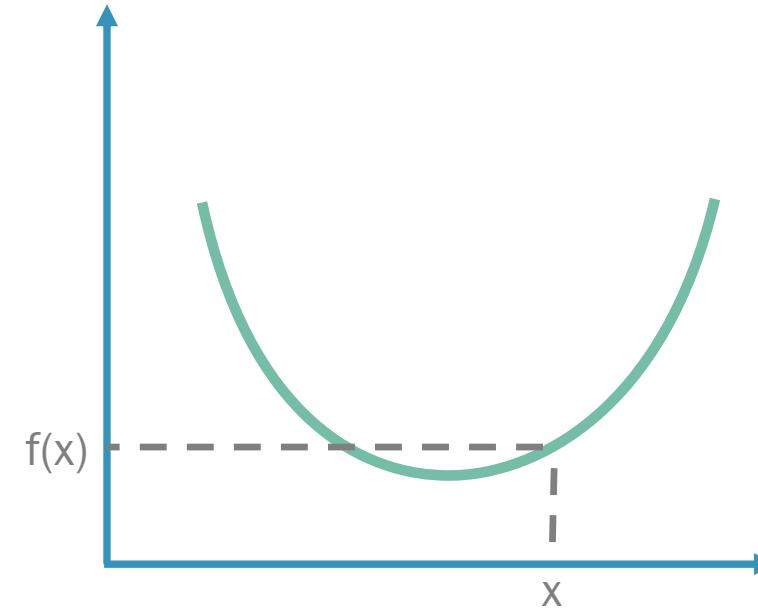
- **Derivatives**
 - $y = f(x)$
 - How a small change in x changes y



It's Calculus Time



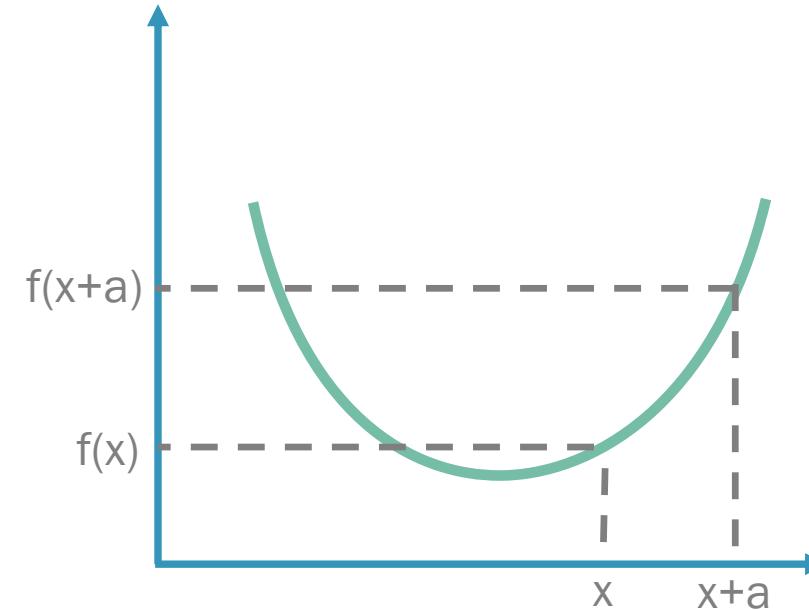
- **Derivatives**
 - $y = f(x)$
 - How a small change in x changes y



It's Calculus Time



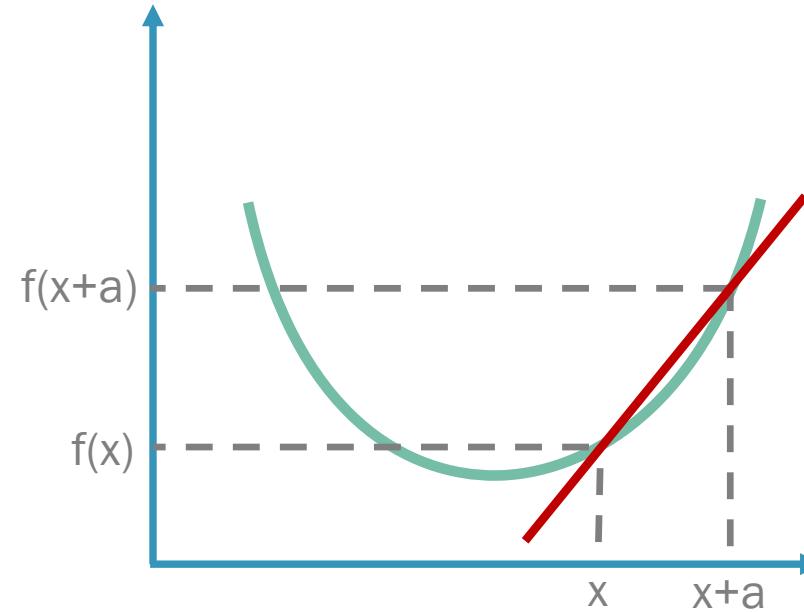
- **Derivatives**
 - $y = f(x)$
 - How a small change in x changes y



It's Calculus Time



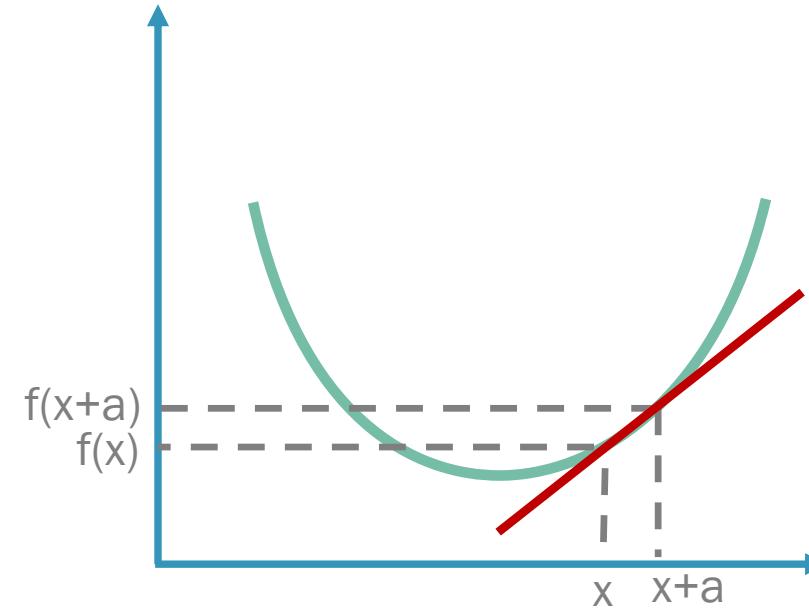
- **Derivatives**
 - $y = f(x)$
 - How a small change in x changes y



It's Calculus Time



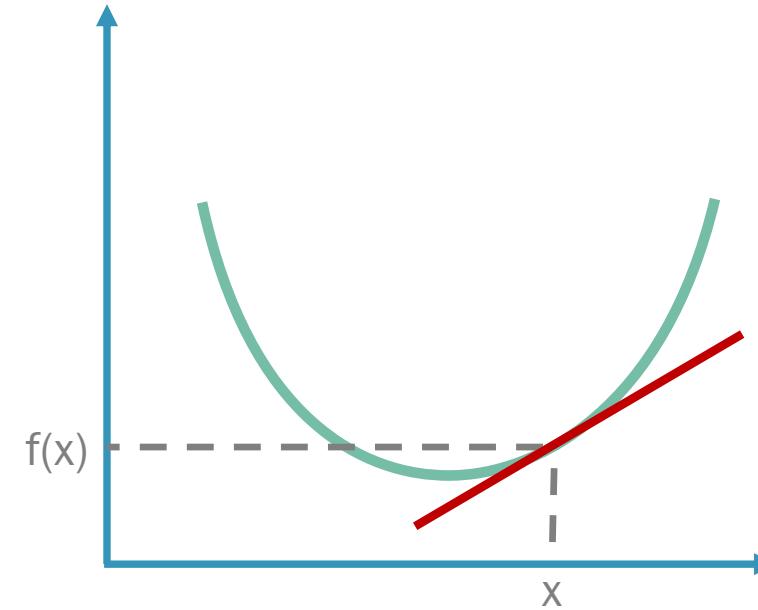
- **Derivatives**
 - $y = f(x)$
 - How a small change in x changes y



It's Calculus Time



- **Derivatives**
 - $y = f(x)$
 - How a small change in x changes y

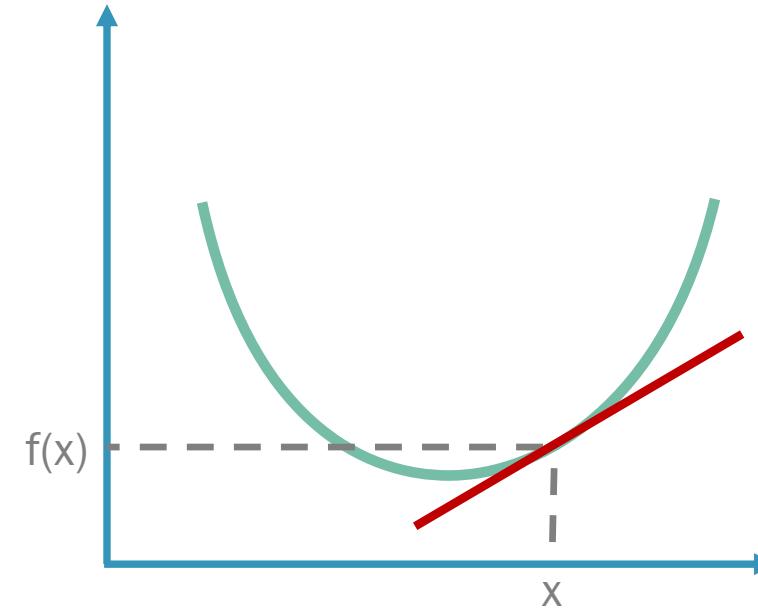


It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{(x+a)-x}$ or $\frac{dy}{dx}$

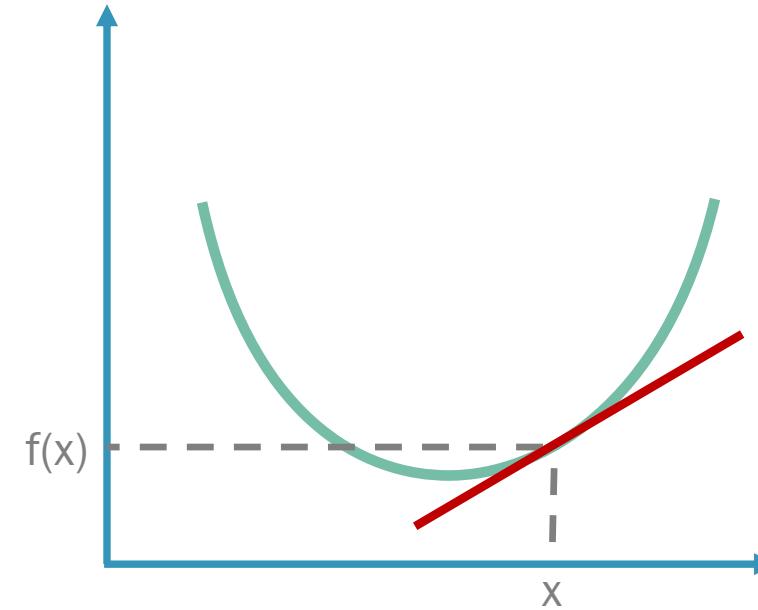


It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$

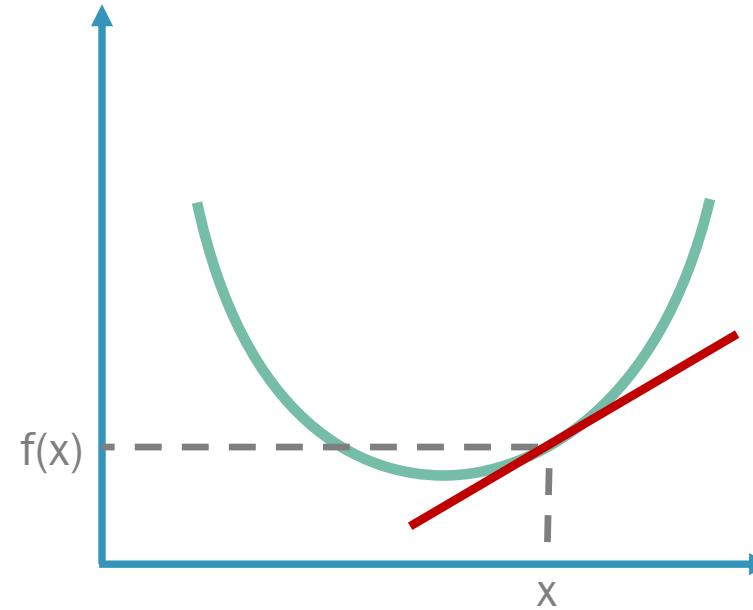


It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- **Slope** of the tangent line



It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

It's Calculus Time

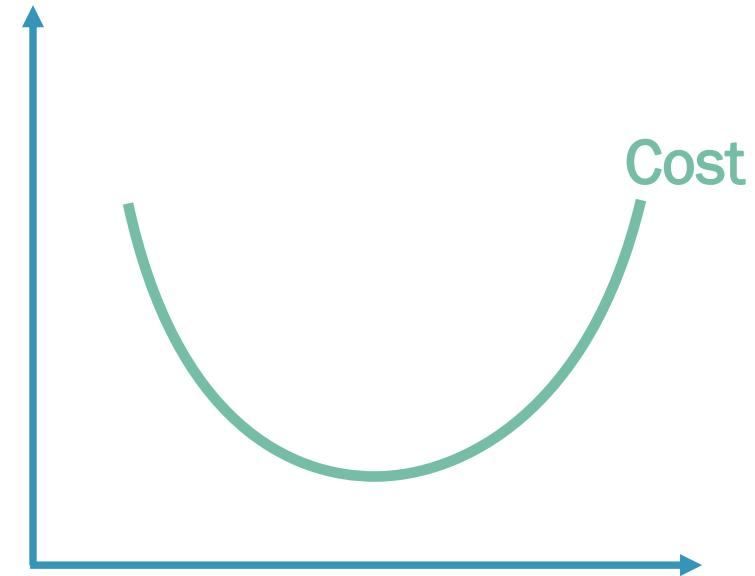


- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



It's Calculus Time

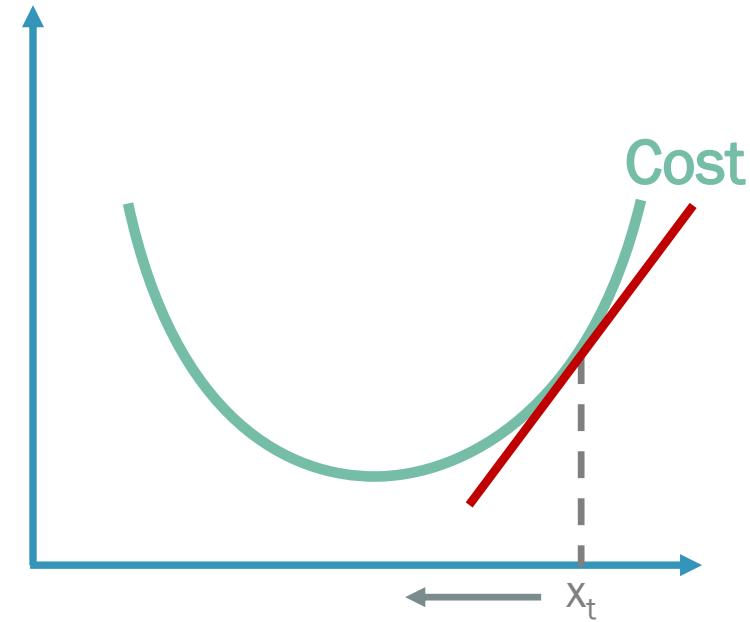


- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



It's Calculus Time

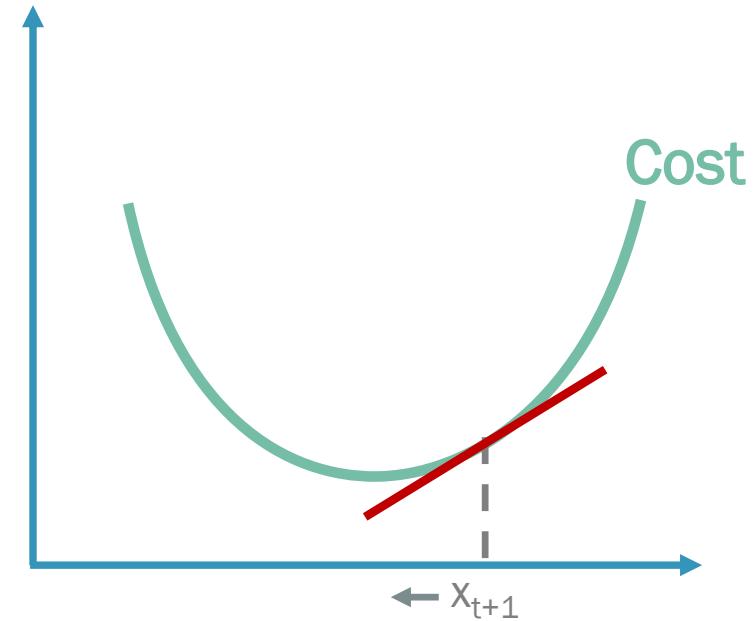


- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



It's Calculus Time

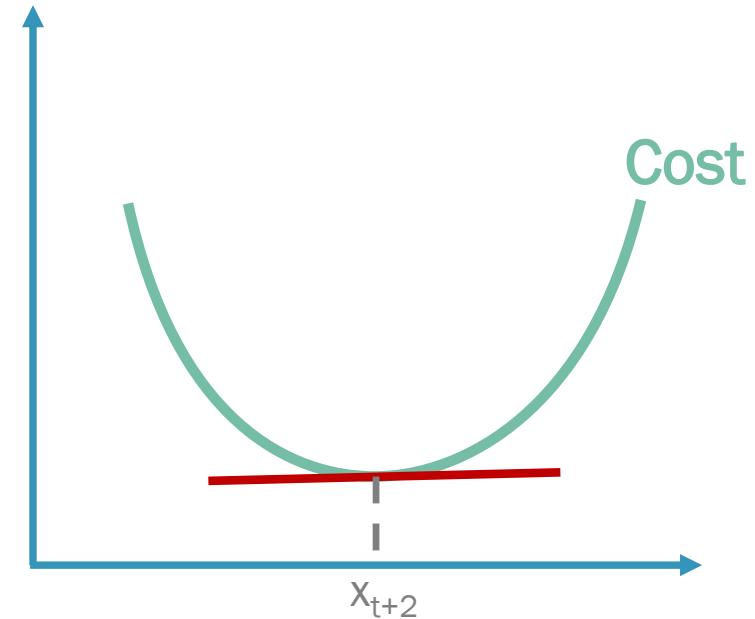


- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



It's Calculus Time

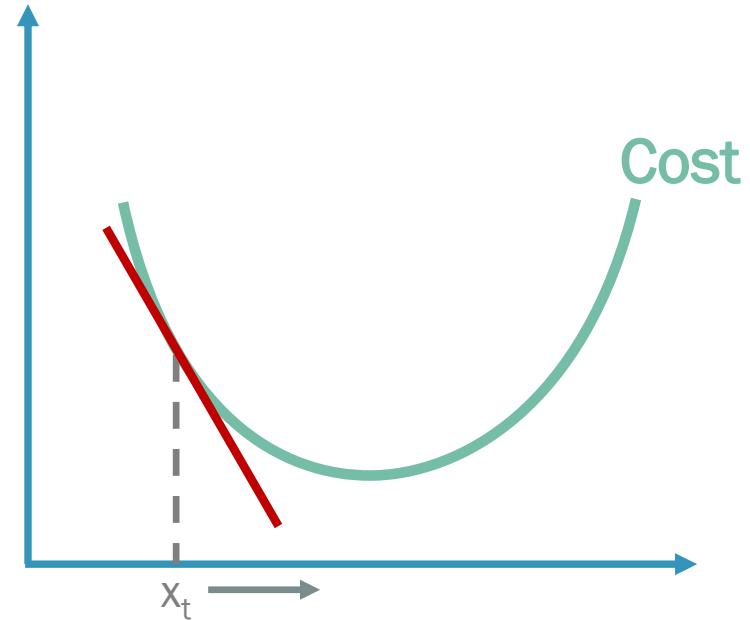


- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

- $y = f(x_1, x_2, \dots, x_k)$

It's Calculus Time



- **Derivatives**

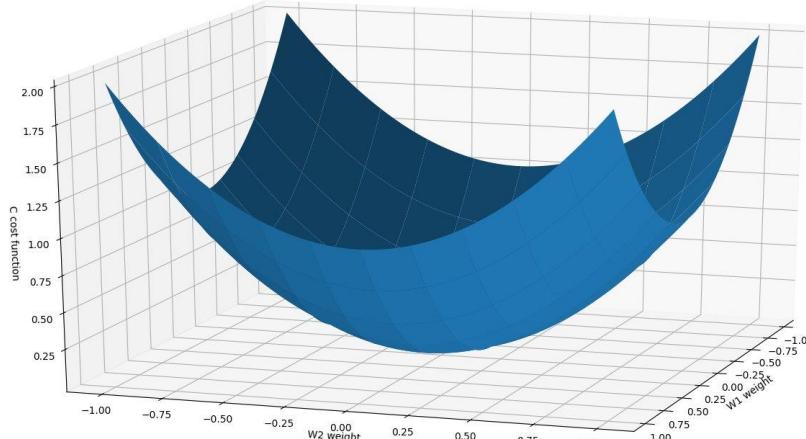
- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

- $y = f(x_1, x_2, \dots, x_k)$



It's Calculus Time



- **Derivatives**

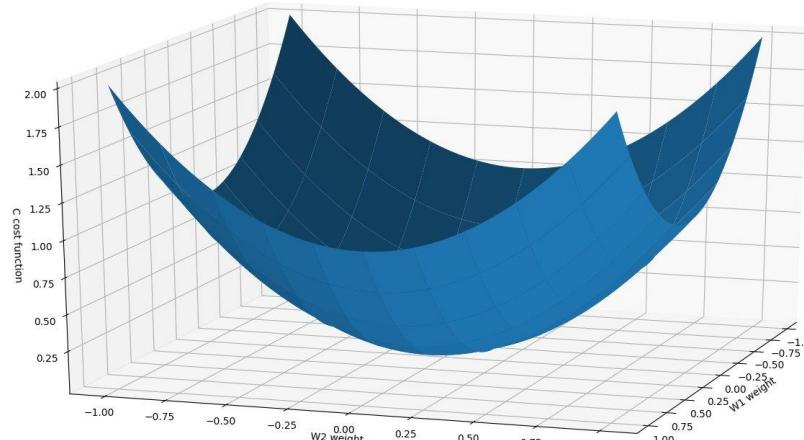
- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

- $y = f(x_1, x_2, \dots, x_k)$
- How a small change in x_i changes y , keeping all $x_{k \neq i}$ constant



It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in x changes y
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$ or $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

- $y = f(x_1, x_2, \dots, x_k)$
- How a small change in x_i changes y , keeping all $x_{k \neq i}$ constant
- $\frac{\partial y}{\partial x}$ instead of $\frac{dy}{dx}$

The chain rule of derivatives

The chain rule of derivatives

$$\frac{d}{dx} g(f(x)) = \frac{dg}{df} \frac{df}{dx}$$

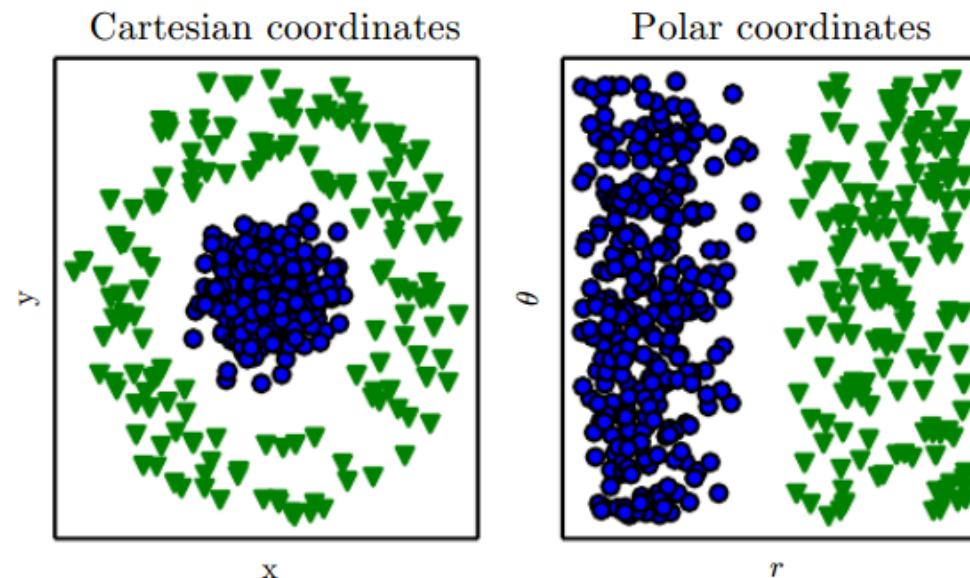
Perceptron

Perceptron

- Linear binary classifier

Perceptron

- Linear binary classifier
 - Linearly separable data



Perceptron

- Linear binary classifier
 - Linearly separable data
 - One out of two classes



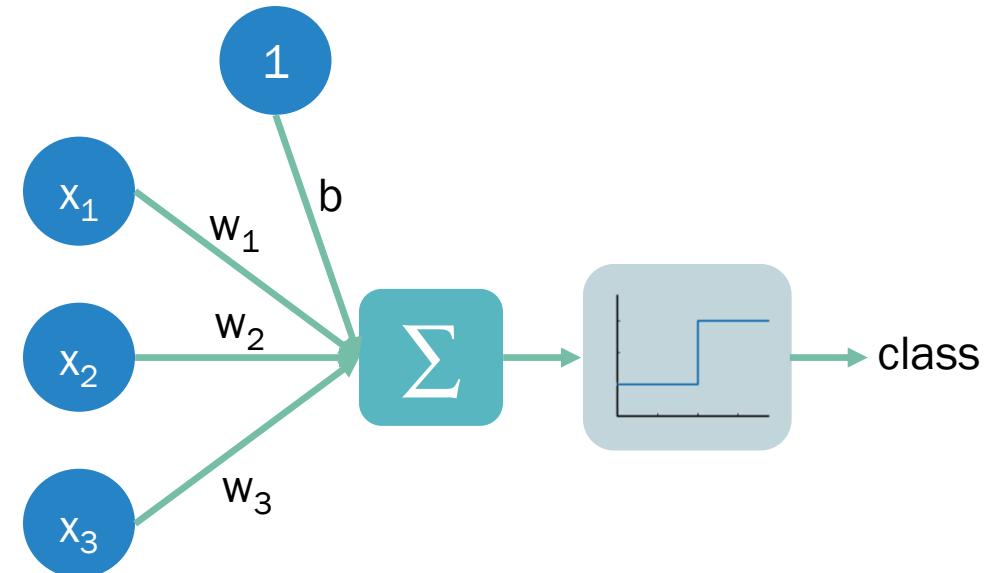
Perceptron

- **Linear binary classifier**
 - Linearly separable data
 - One out of two classes
 - Supervised learning



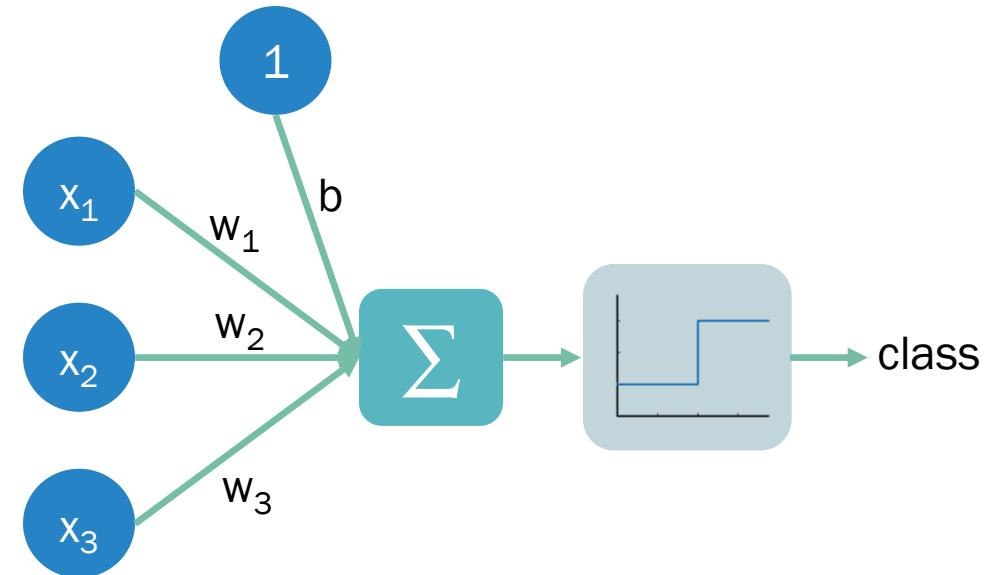
Perceptron

- Linear binary classifier
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- Components



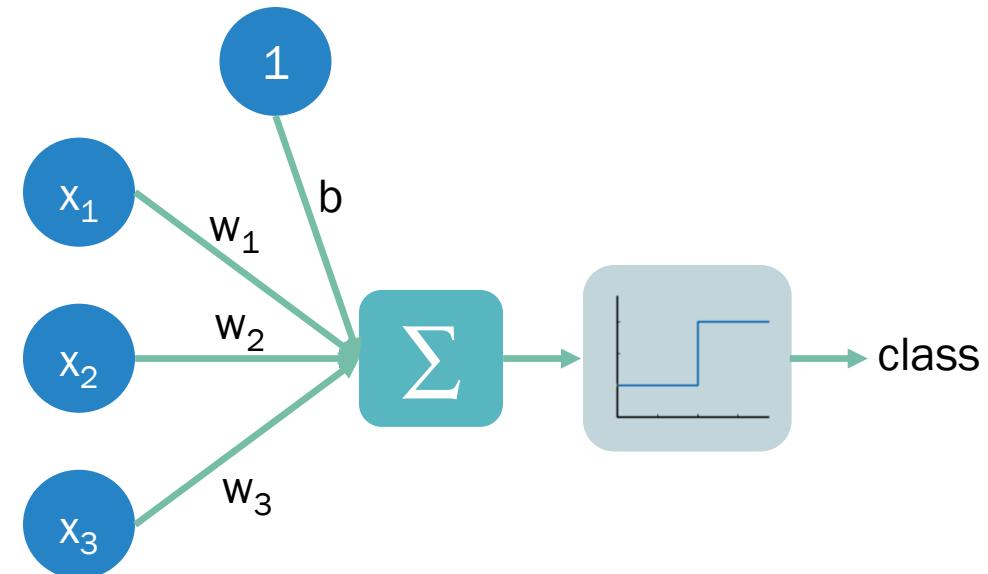
Perceptron

- Linear binary classifier
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- Components
 - Input values



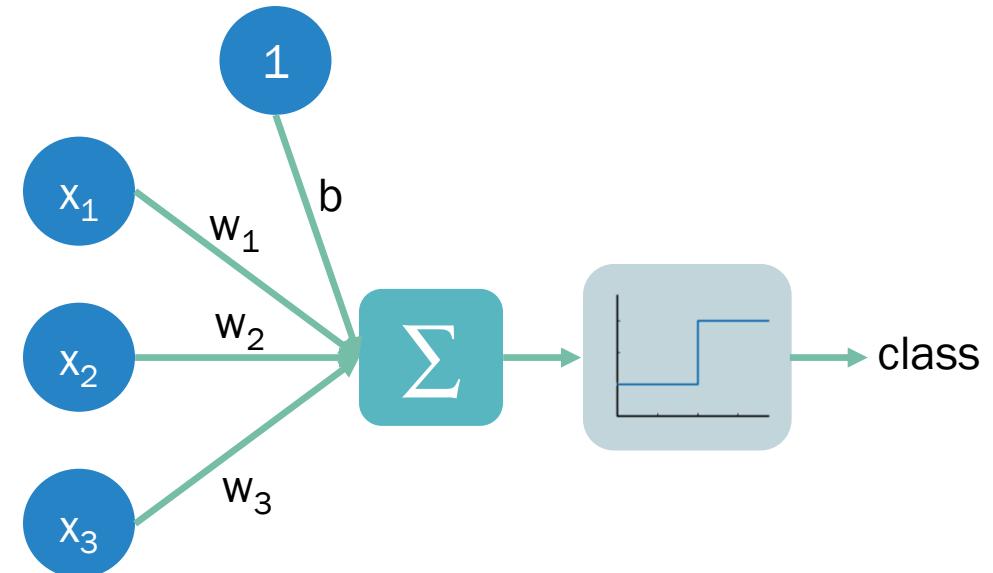
Perceptron

- **Linear binary classifier**
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- **Components**
 - Input values
 - Weights and Biases



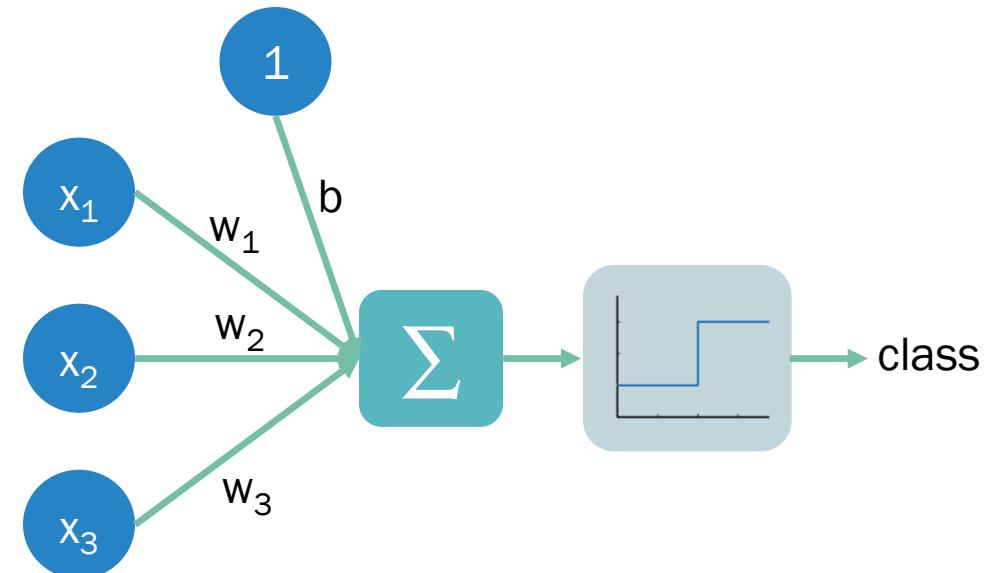
Perceptron

- **Linear binary classifier**
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- **Components**
 - Input values
 - Weights and Biases
 - Sum



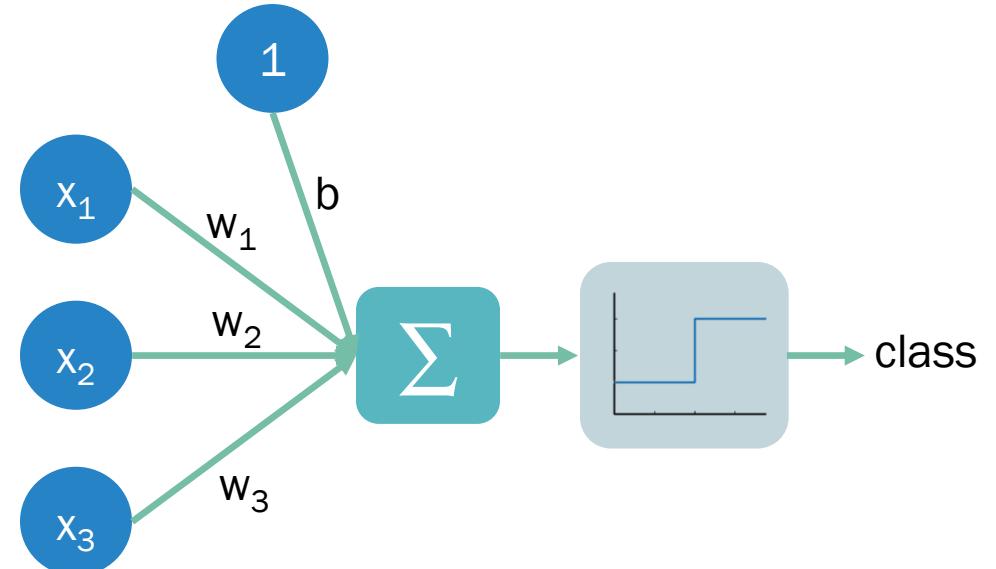
Perceptron

- **Linear binary classifier**
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- **Components**
 - Input values
 - Weights and Biases
 - Sum
 - Activation function



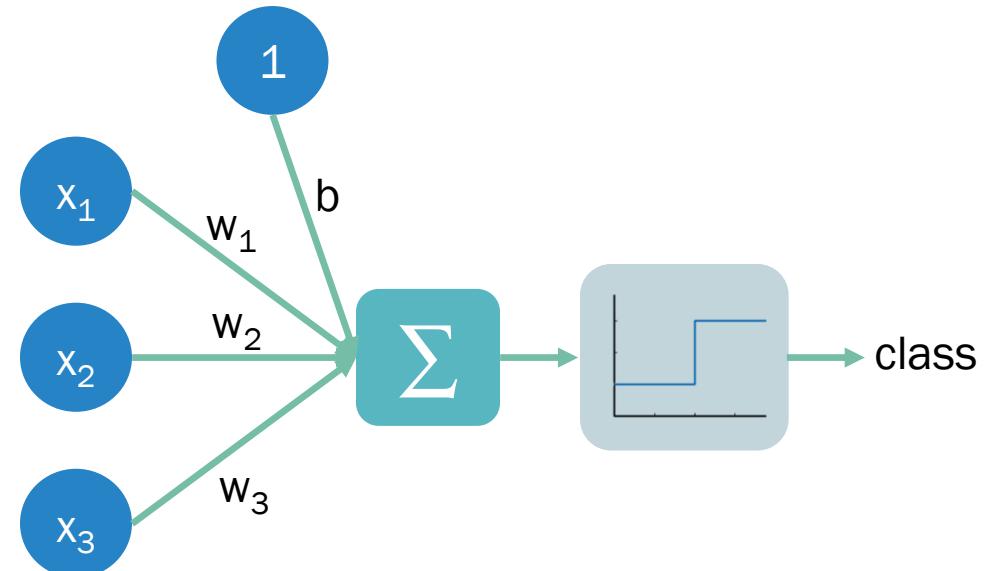
Perceptron

- **Linear binary classifier**
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- **Components**
 - Input values
 - Weights and Biases
 - Sum
 - Activation function
 - Weight update from errors (*training*)



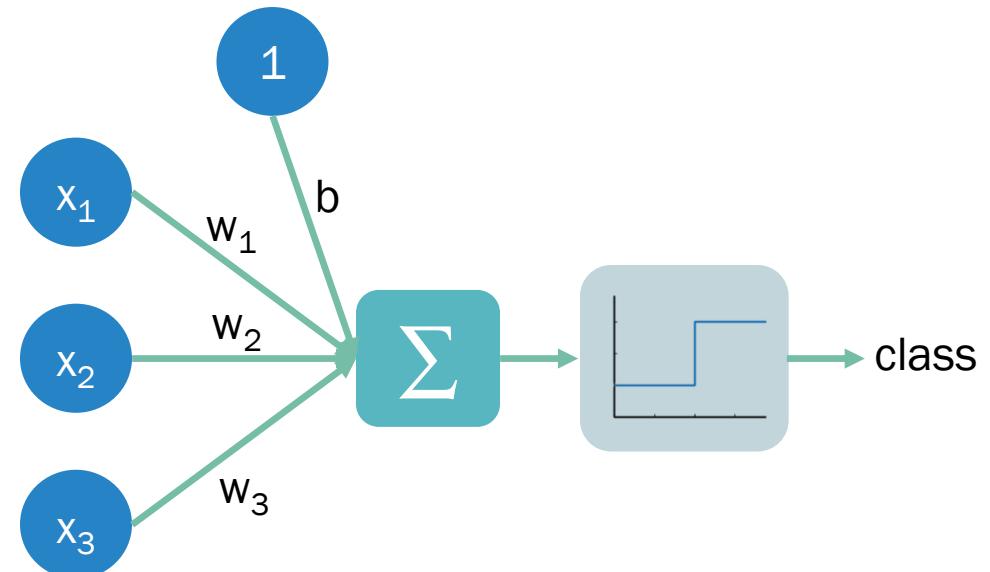
Perceptron

- **Linear binary classifier**
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- **Components**
 - Input values
 - Weights and Biases
 - Sum
 - Activation function
 - Weight update from errors (*training*)
- **Steps**



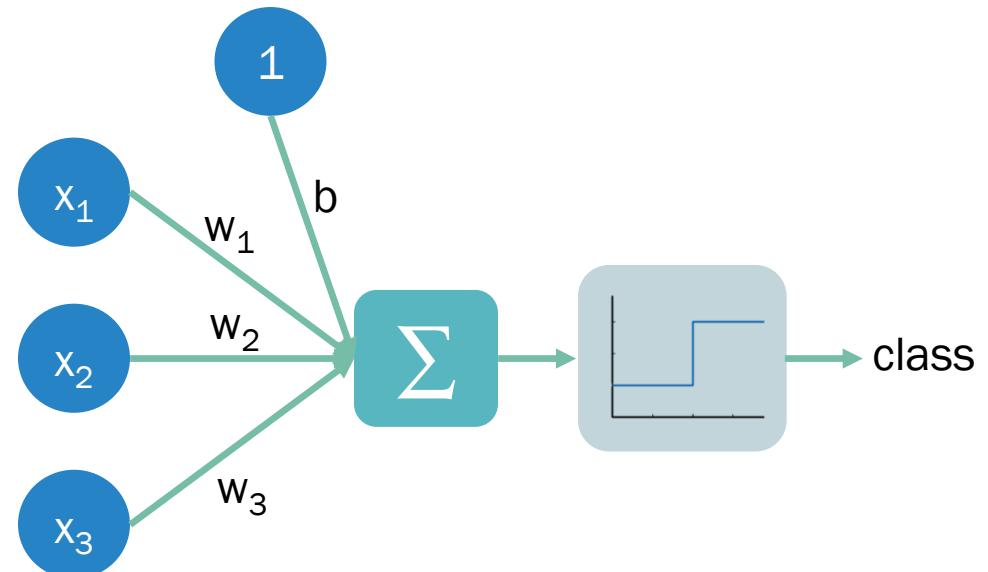
Perceptron

- Linear binary classifier
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- Components
 - Input values
 - Weights and Biases
 - Sum
 - Activation function
 - Weight update from errors (*training*)
- Steps
 1. Multiply inputs by the weights



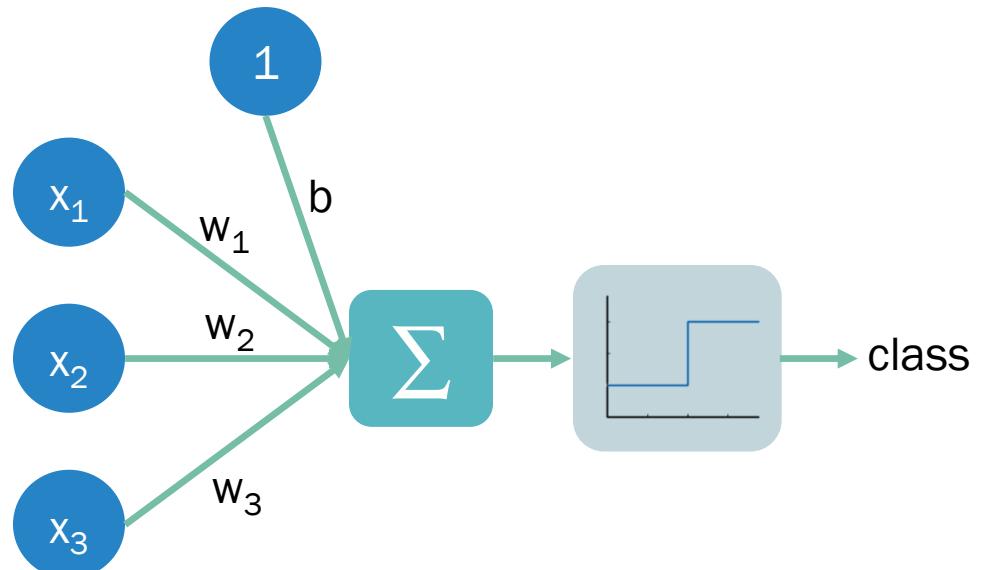
Perceptron

- **Linear binary classifier**
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- **Components**
 - Input values
 - Weights and Biases
 - Sum
 - Activation function
 - Weight update from errors (*training*)
- **Steps**
 1. Multiply inputs by the weights
 2. Add up the values



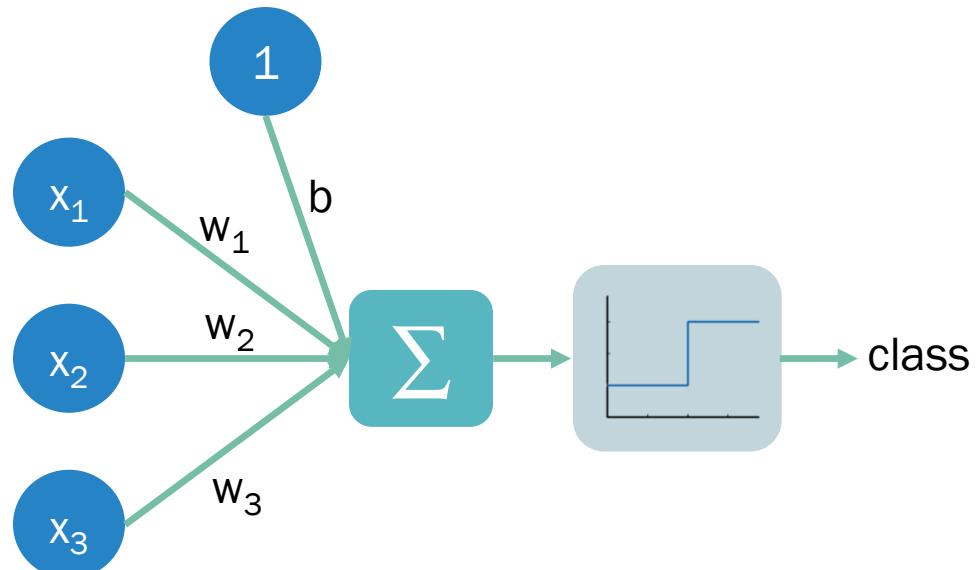
Perceptron

- **Linear binary classifier**
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- **Components**
 - Input values
 - Weights and Biases
 - Sum
 - Activation function
 - Weight update from errors (*training*)
- **Steps**
 1. Multiply inputs by the weights
 2. Add up the values
 3. Apply the activation function



Perceptron

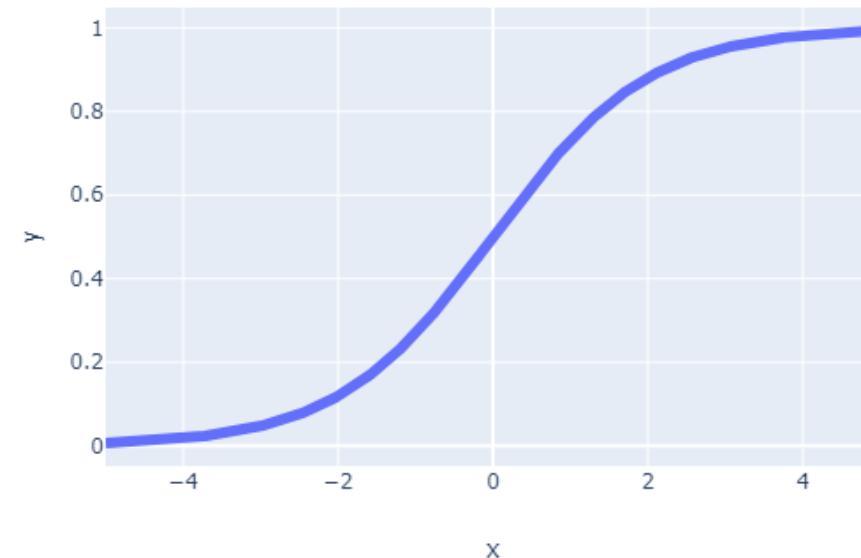
- Linear binary classifier
 - Linearly separable data
 - One out of two classes
 - Supervised learning
- Components
 - Input values
 - Weights and Biases
 - Sum
 - Activation function
 - Weight update from errors (*training*)
- Steps
 1. Multiply inputs by the weights
 2. Add up the values
 3. Apply the activation function
 4. Update the weights given error (*training*)



Step & Error functions

Step & Error functions

- Sigmoid function

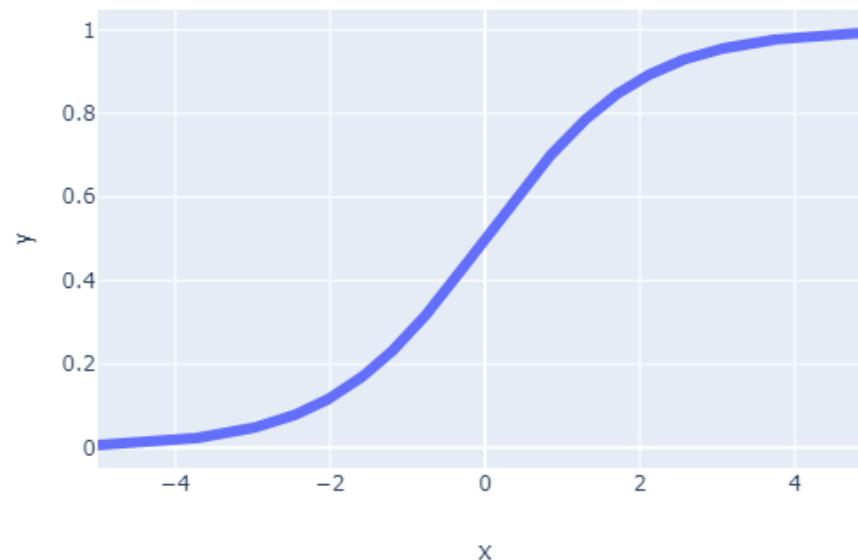


Step & Error functions

- **Sigmoid function**

- $\frac{1}{1+e^{-x}}$

- Can be interpreted as probability



Step & Error functions

- Sigmoid function
 - $\frac{1}{1+e^{-x}}$
 - Can be interpreted as probability
- Cross-entropy



Step & Error functions

- **Sigmoid function**

- $\frac{1}{1+e^{-x}}$

- Can be interpreted as probability

- **Cross-entropy**

- $-\sum_i(t_i \ln(y_i) + (1 - t_i) \ln(1 - y_i))$
- Distance between two distributions



Step & Error functions

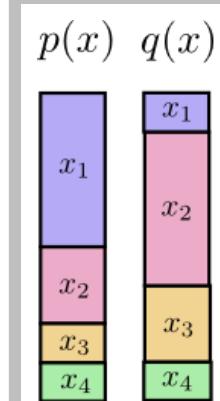
- **Sigmoid function**

- $\frac{1}{1+e^{-x}}$

- Can be interpreted as probability

- **Cross-entropy**

- $-\sum_i(t_i \ln(y_i) + (1 - t_i) \ln(1 - y_i))$
- Distance between two distributions
- Rooted on information theory
 - Huffman coding



Cross-Entropy: $H_p(q)$

Average Length
of message from $q(x)$
using code for $p(x)$.

* From [Visual Information Theory](#) by Christopher Olah

The Internals

Is it all a black box?

"For every action there is an equal and opposite reaction "

Newton's 3rd Law of Motion

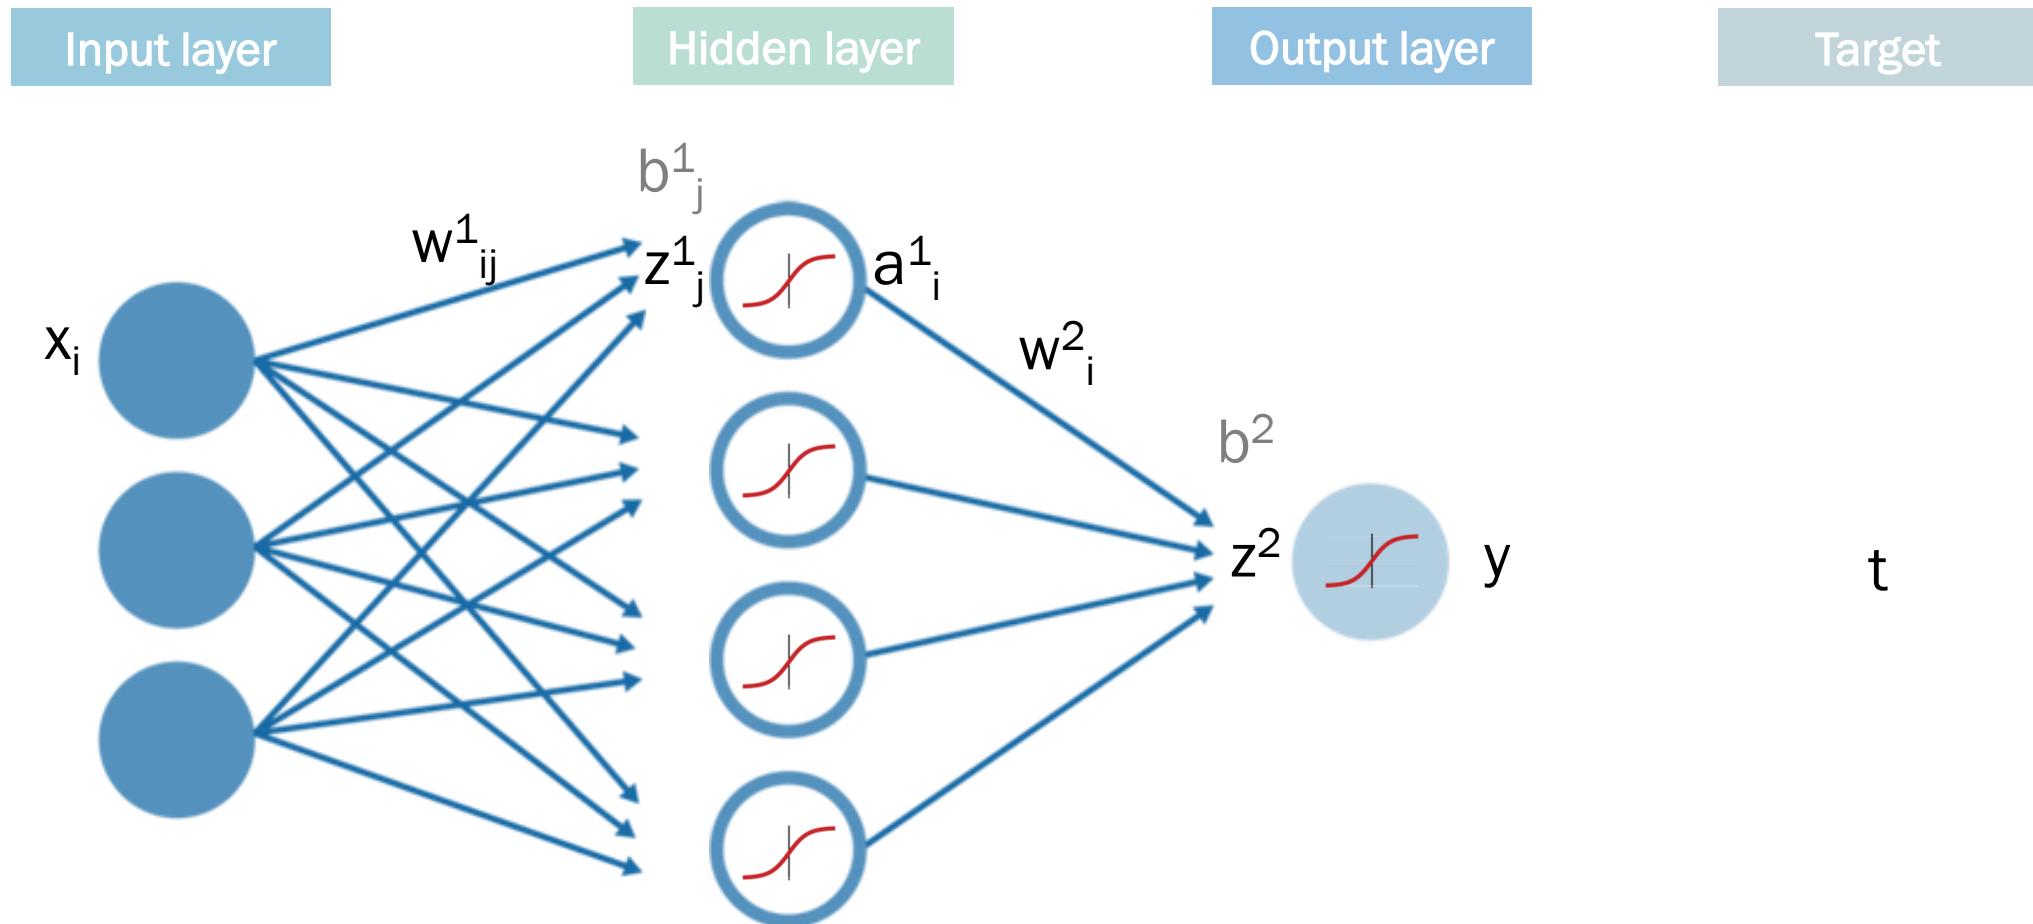
Setup

Single datum

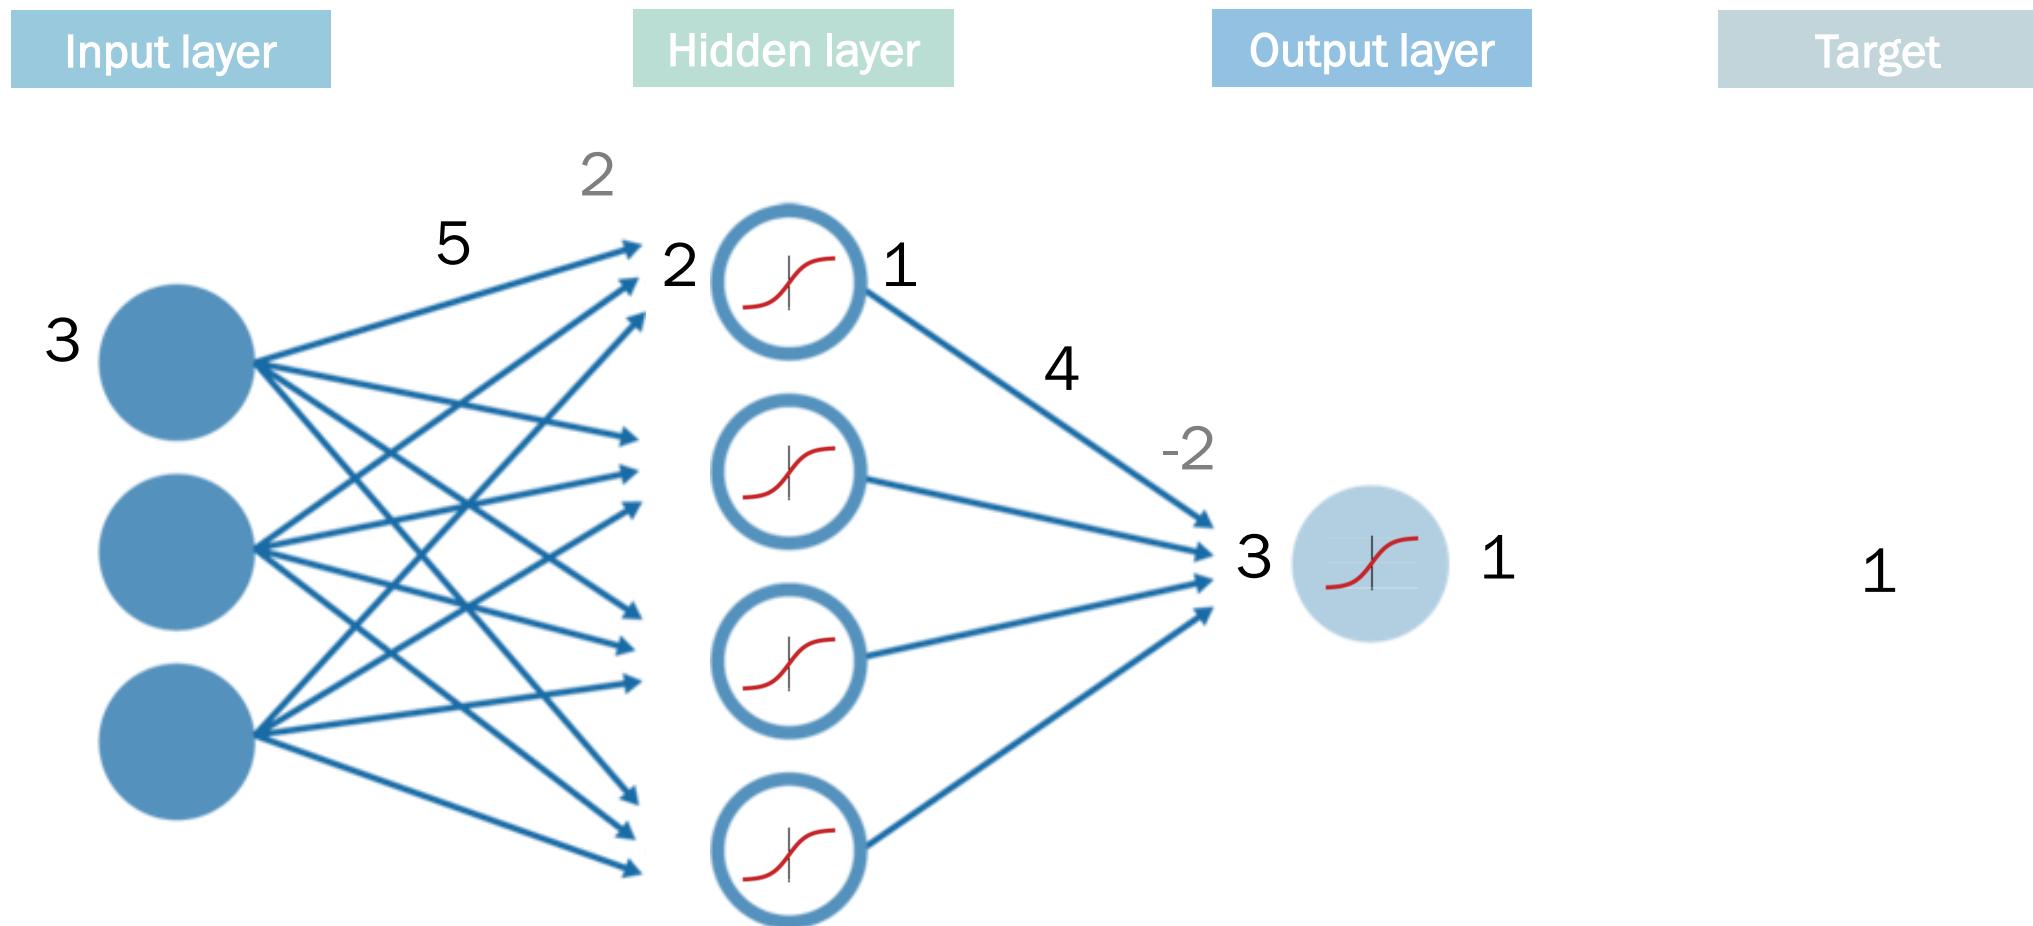
Binary classification
problem

Objective:
 $\arg \min_{w,b} \text{Error}$

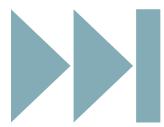
Anatomy of a neural network



Anatomy of a neural network



Derivation Steps



Forward
Propagation



Error
Computation



Back
Propagation



Parameters
update

Implementation

Does it even work?

"Machines take me by surprise with great frequency "

Alan Turing

Just a tiny bit more of fundamentals

- Dot product of two matrices

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`?
```



Just a tiny bit more of fundamentals



- Dot product of two matrices
 - Iterating over the arrays
 - 247 ms

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c_slow = np.zeros((100, 100))

for i in range(a.shape[0]):
    for j in range(b.shape[1]):
        acc = 0
        for k in range(a.shape[1]):
            acc += a[i, k] * b[k, j]
        c_slow[i, j] = acc
```

Just a tiny bit more of fundamentals



- Dot product of two matrices
 - Iterating over the arrays
 - 247 ms
 - Numpy's dot product
 - 110 µs

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

Just a tiny bit more of fundamentals



- Dot product of two matrices
 - Iterating over the arrays
 - 247 ms
 - Numpy's dot product
 - 110 µs
- **Vectorization** – rewriting loops into efficient parallelizable operations

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

Just a tiny bit more of fundamentals



- Dot product of two matrices
 - Iterating over the arrays
 - 247 ms
 - Numpy's dot product
 - 110 µs
- **Vectorization** – rewriting loops into efficient parallelizable operations
 - “Closer-to-the-metal” computations

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

Just a tiny bit more of fundamentals



- Dot product of two matrices
 - Iterating over the arrays
 - 247 ms
 - Numpy's dot product
 - 110 µs
- **Vectorization** – rewriting loops into efficient parallelizable operations
 - “Closer-to-the-metal” computations
 - Prone to parallelization

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

Just a tiny bit more of fundamentals



- Dot product of two matrices
 - Iterating over the arrays
 - 247 ms
 - Numpy's dot product
 - 110 µs
- **Vectorization** – rewriting loops into efficient parallelizable operations
 - “Closer-to-the-metal” computations
 - Prone to parallelization
 - Cleaner code

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

ENOUGH TALK



LET'S CODE

memegenerator.net

<http://bit.ly/2GFMJJP>



Today we
learned

Today we learned

What we did:

- Derived the math behind traditional neural networks for binary classification
- Implemented and evaluated them using Numpy

Today we learned

What we did:

- Derived the math behind traditional neural networks for binary classification
- Implemented and evaluated them using Numpy

Potential paths of exploration:

- Regression
 - Linear output
 - Mean squared error
- Multiclass
 - Softmax activation
 - Cross entropy
- Refactor the implementation
 - Easily add layers
 - ...

Bibliography

- [Deep Learning](#) by Goodfellow et al. for images and inspiration
- [Notes on Backpropagation](#)
- [Census income dataset](#) from [UCI Repository](#)
- [Visual Information Theory](#) by [Christopher Olah](#)
- [Neural Networks MOOC](#) by [Geoffrey Hinton](#)
- [What the Hell is a Perceptron](#) by [Sagar Sharma](#)

Acknowledgements

- Icons made by [Freepik](#), [dDara](#), [surang](#), [Vectors Market](#), obtained from [Flaticon](#)

Thank you



[@diogojapinto](https://twitter.com/diogojapinto)



github.com/diogojapinto



pt.linkedin.com/in/diogojapinto