


# Deep Dissection

Understanding the ins and outs of Neural Networks

*by Diogo Pinto*



The road  
ahead

# The road ahead

## Motivation

- Why deep learning
- Playground

# The road ahead

## Motivation

- Why deep learning
- Playground

## Fundamentals

- Derivatives & Chain rule
- Perceptron
- Sigmoid function & Cross-entropy

# The road ahead

## Motivation

- Why deep learning
- Playground

## Fundamentals

- Derivatives & Chain rule
- Perceptron
- Sigmoid function & Cross-entropy

## The Internals

- Forward propagation
- Error computation
- Back propagation
- Parameters update

# The road ahead

## Motivation

- Why deep learning
- Playground

## Fundamentals

- Derivatives & Chain rule
- Perceptron
- Sigmoid function & Cross-entropy

## The Internals

- Forward propagation
- Error computation
- Back propagation
- Parameters update

## Coffee Break



# The road ahead

## Motivation

- Why deep learning
- Playground

## Fundamentals

- Derivatives & Chain rule
- Perceptron
- Sigmoid function & Cross-entropy

## The Internals

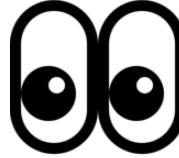
- Forward propagation
- Error computation
- Back propagation
- Parameters update

## Coffee Break



## Implementation

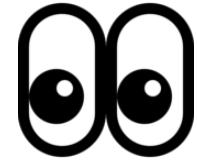
- Binary Classification – Census Income dataset

Who are you people? 

# Who are you people?


- Who is used to work with Python?

# Who are you people?




- Who is used to work with Python?
- Who already did something **close to this**?

# Who are you people?




- Who is used to work with Python?
- Who already did something close to this?
- Who likes **matrices and letters** in the place of numbers?

# Who are you people?



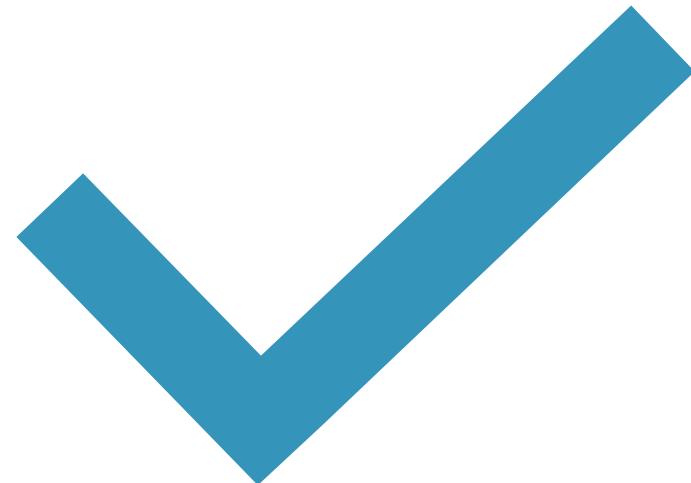
- Who is used to work with Python?
- Who already did something close to this?
- Who likes matrices and letters in the place of numbers?
- Who had a good night of **sleep**?

# Assumptions




# Assumptions

- A bit of **Algebra** and **Calculus**




# Assumptions

- A bit of **Algebra and Calculus**
- Introductory level of **Machine Learning** knowledge



# Assumptions

- A bit of **Algebra** and **Calculus**
- Introductory level of **Machine Learning** knowledge



Difficulty can be calibrated, give feedback!

# Motivation

---

Why am I here?

---

*"Don't limit your challenges, challenge your limits"*

*Unknown*

The trick is in the representation...

# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive

# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive

01100  
10110  
11110

# The trick is in the representation...


- Machine learning is traditionally **feature engineering** intensive

01100  
10110  
11110



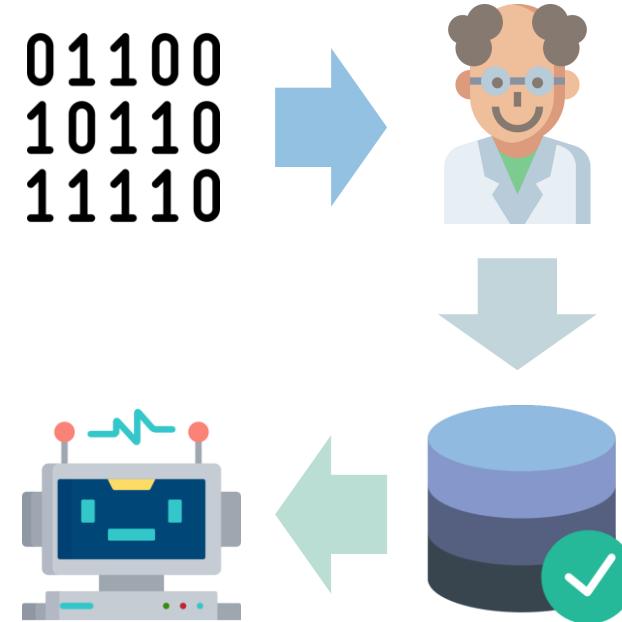
# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive



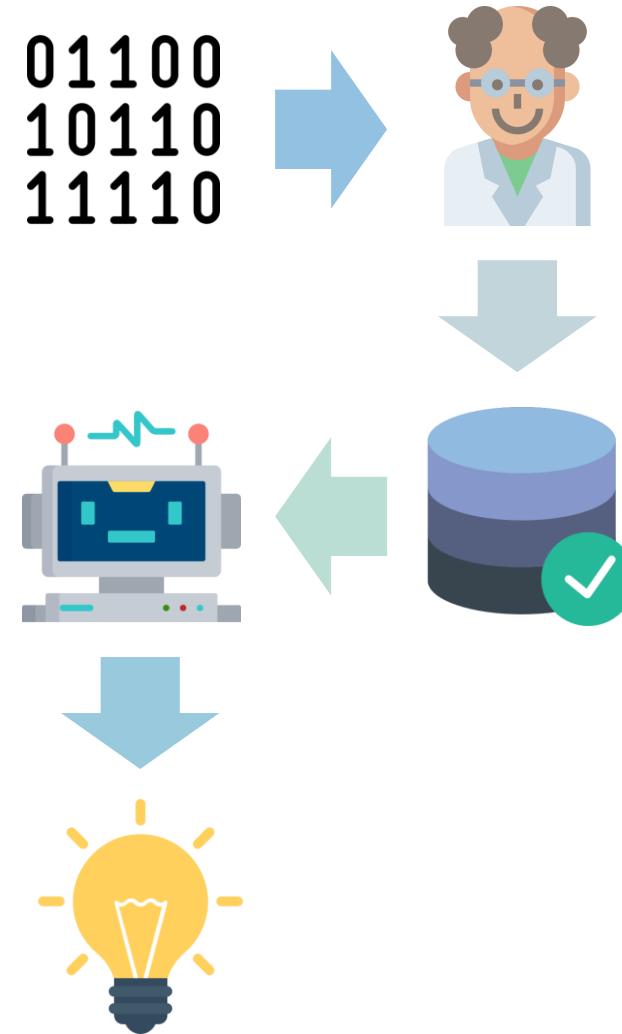
# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive



# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive




# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**

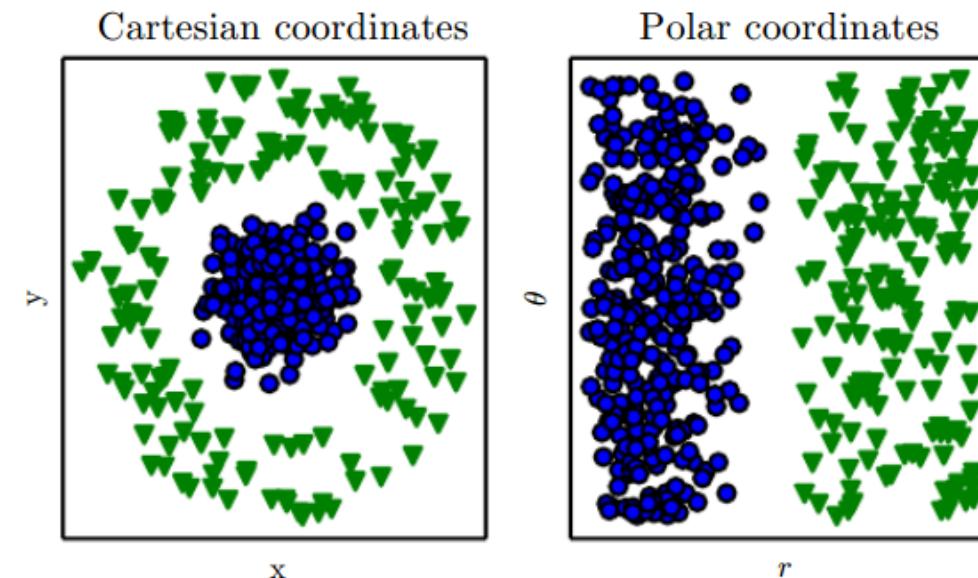
# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**




# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**




# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**
  - Disentangle them



# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**
  - Disentangle them
  - Discard op the unnecessary ones

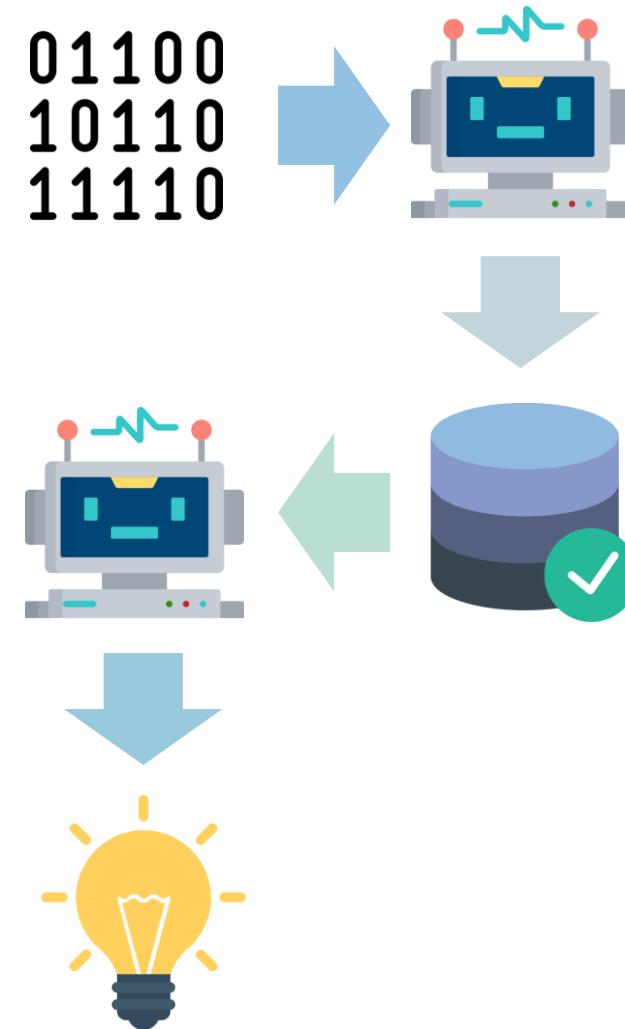


# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**
  - Disentangle them
  - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process

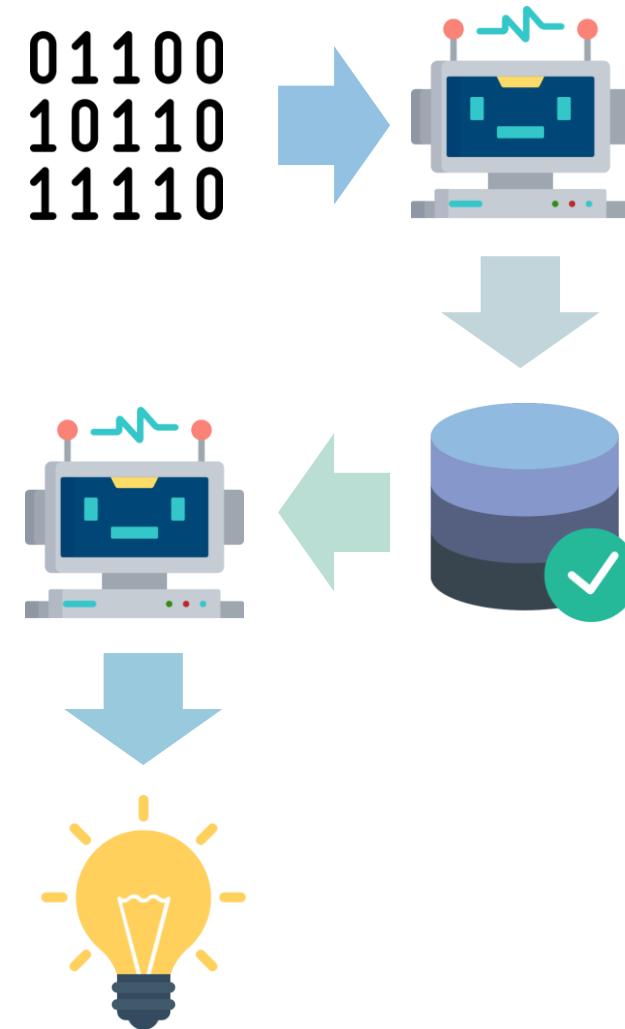
# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**
  - Disentangle them
  - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process



# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**
  - Disentangle them
  - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process
  - Learning **high-level concepts** can be as hard as the original problem

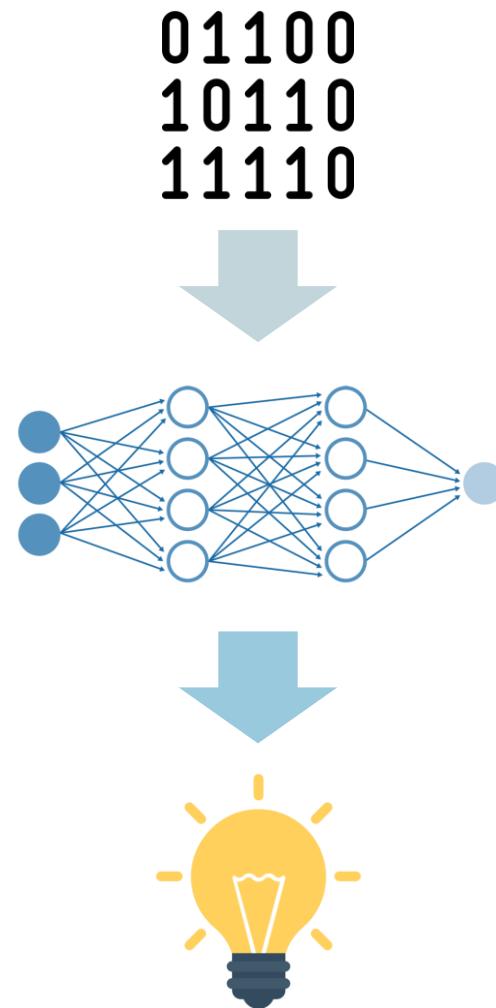


# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**
  - Disentangle them
  - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process
  - Learning **high-level concepts** can be as hard as the original problem
- **Deep learning** tackles this heads on


# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**
  - Disentangle them
  - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process
  - Learning **high-level concepts** can be as hard as the original problem
- **Deep learning** tackles this heads on



# The trick is in the representation...

- Machine learning is traditionally **feature engineering** intensive
  - Identifying the **sources of influence**
  - Disentangle them
  - Discard op the unnecessary ones
- **Representation learning** seeks to automate this process
  - Learning **high-level concepts** can be as hard as the original problem
- **Deep learning** tackles this heads on
  - **Hierarchical combination** of simpler concepts into more complex ones



# A quick visual demo...

[Neural Networks Playground](#)

# Fundamentals

---

What do I need to know?

---

*"You can't build a great building on a weak foundation."*

*Gordon B. Hinckley*

It's Calculus Time



# It's Calculus Time



- Derivatives

# It's Calculus Time



- Derivatives
  - $y = f(x)$

# It's Calculus Time




- **Derivatives**
  - $y = f(x)$
  - How a small change in  $x$  changes  $y$

# It's Calculus Time



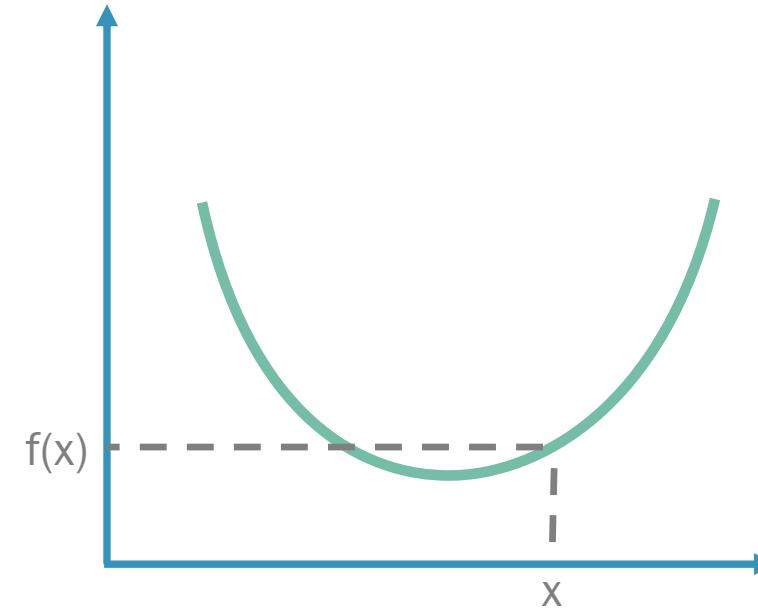
- **Derivatives**
  - $y = f(x)$
  - How a small change in  $x$  changes  $y$



# It's Calculus Time



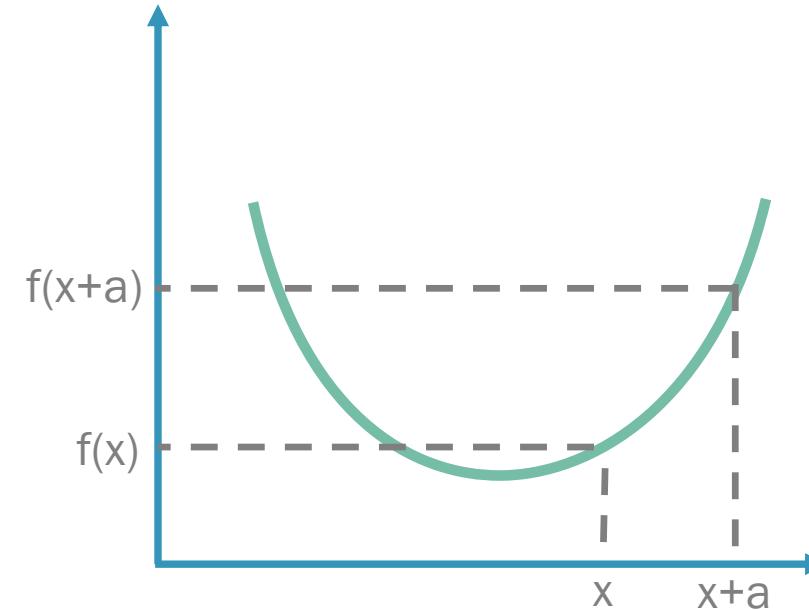
- **Derivatives**
  - $y = f(x)$
  - How a small change in  $x$  changes  $y$



# It's Calculus Time



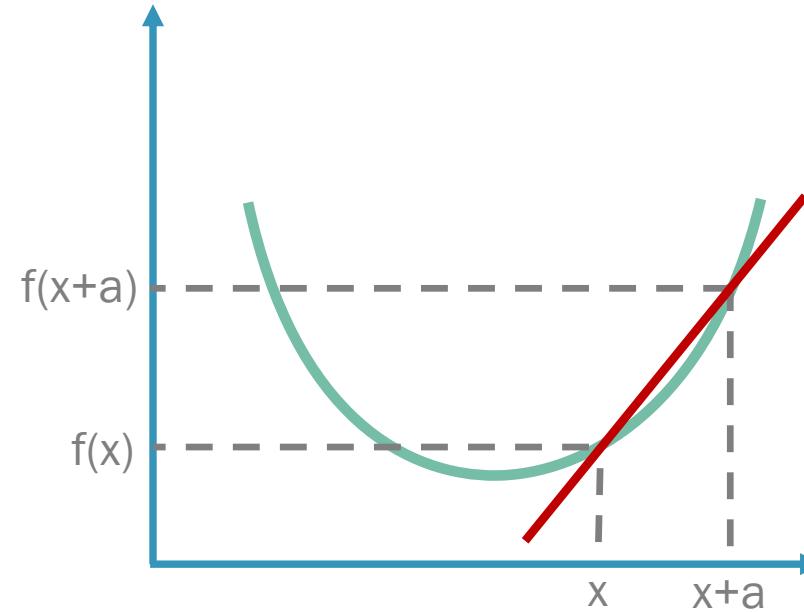
- **Derivatives**
  - $y = f(x)$
  - How a small change in  $x$  changes  $y$



# It's Calculus Time



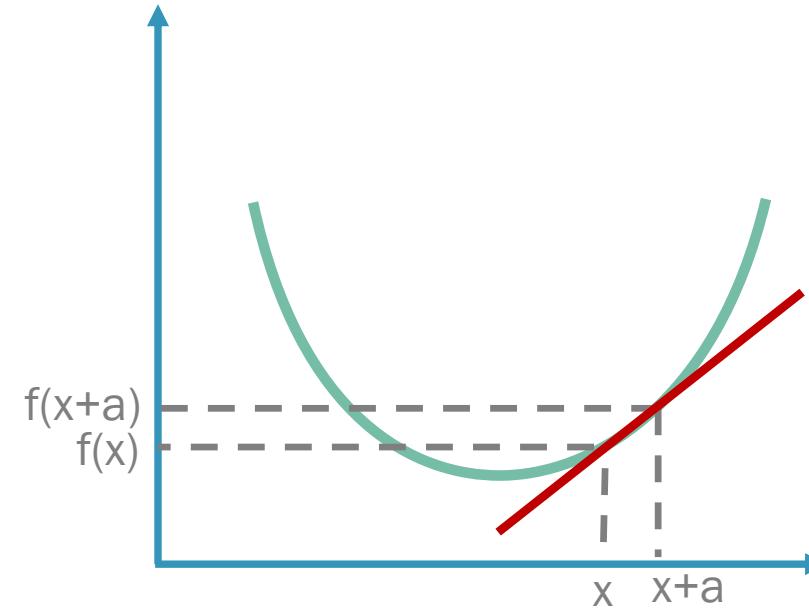
- **Derivatives**
  - $y = f(x)$
  - How a small change in  $x$  changes  $y$



# It's Calculus Time



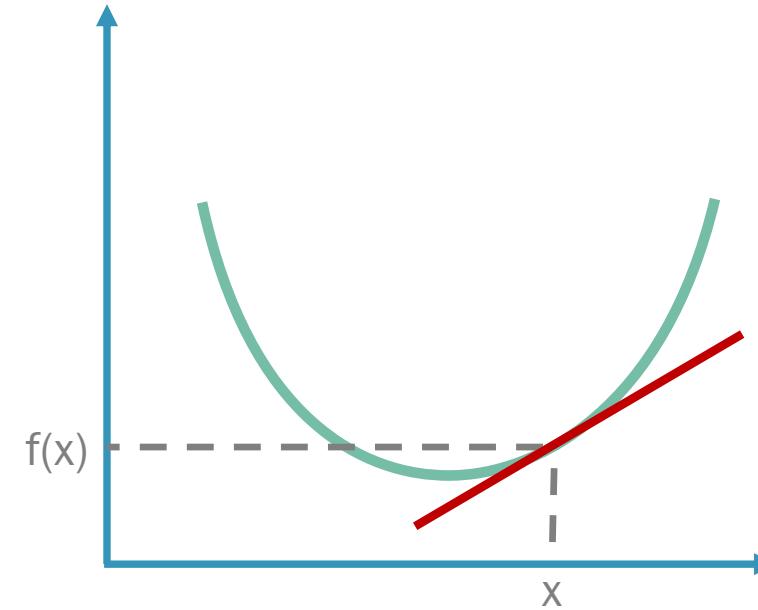
- **Derivatives**
  - $y = f(x)$
  - How a small change in  $x$  changes  $y$



# It's Calculus Time



- **Derivatives**
  - $y = f(x)$
  - How a small change in  $x$  changes  $y$




# It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{(x+a)-x}$  or  $\frac{dy}{dx}$




# It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$

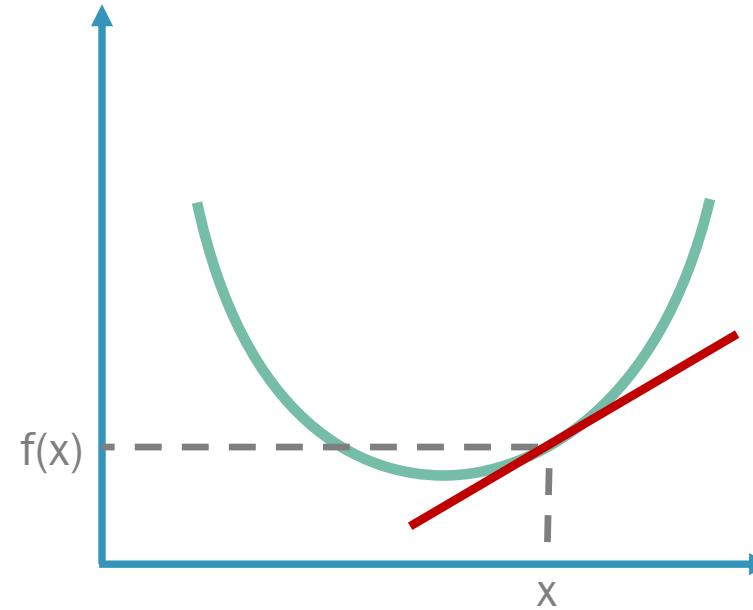


# It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- **Slope** of the tangent line



# It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

# It's Calculus Time

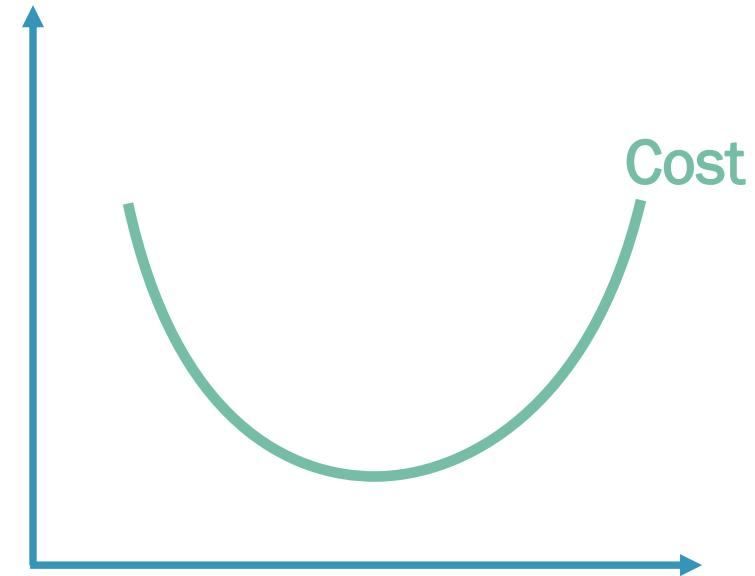


- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



# It's Calculus Time

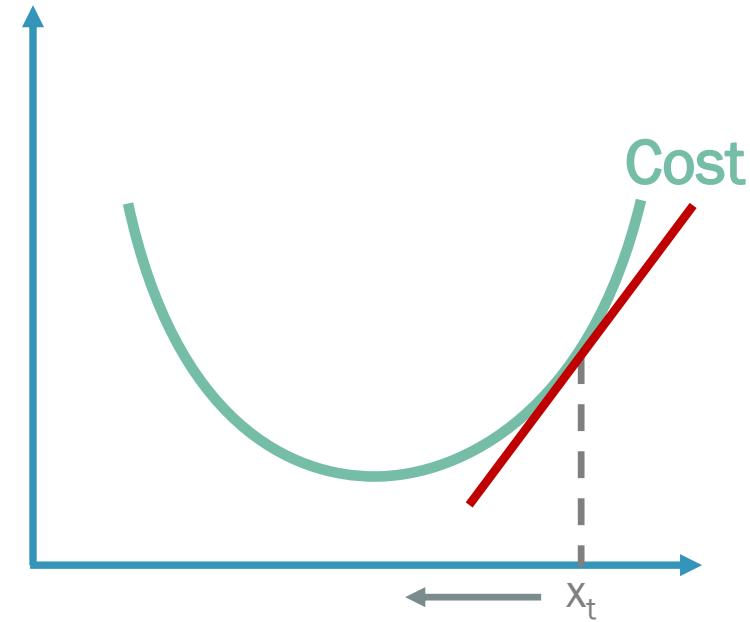


- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



# It's Calculus Time

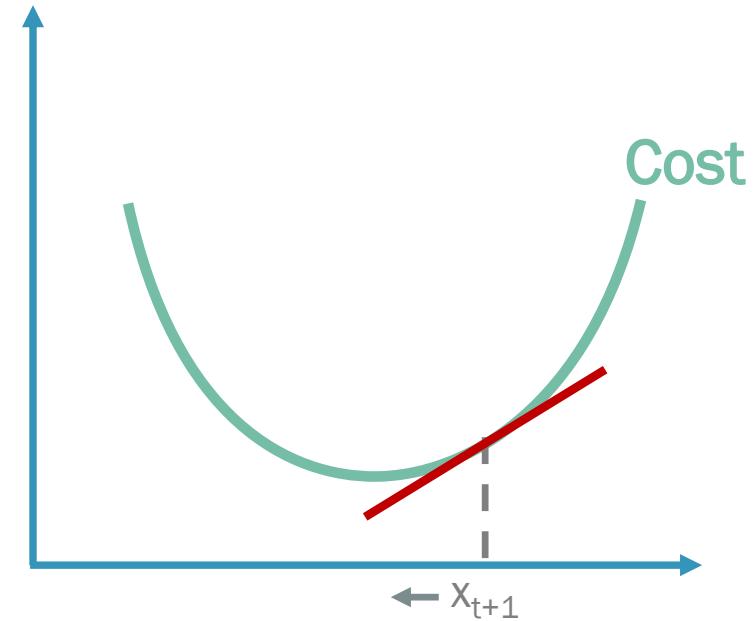


- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



# It's Calculus Time

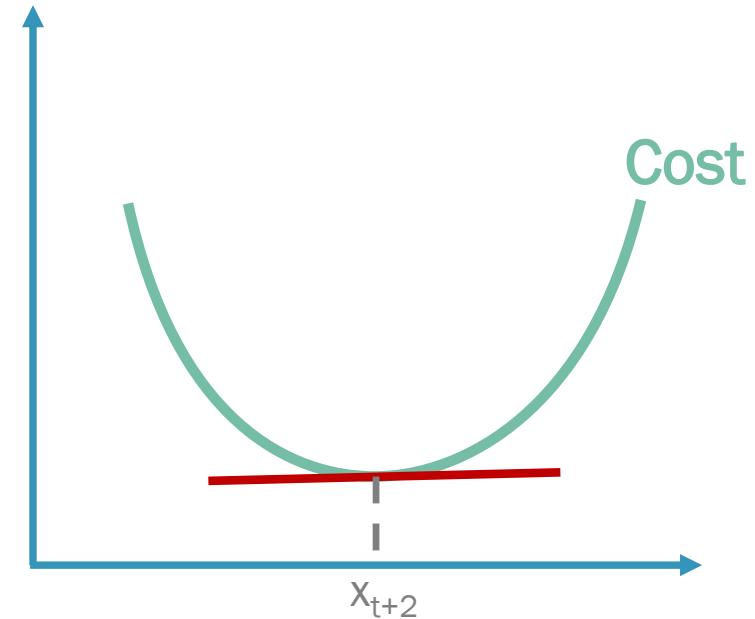


- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



# It's Calculus Time




- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by **subtracting derivative** at the point



# It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

# It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

- $y = f(x_1, x_2, \dots, x_k)$

# It's Calculus Time



- **Derivatives**

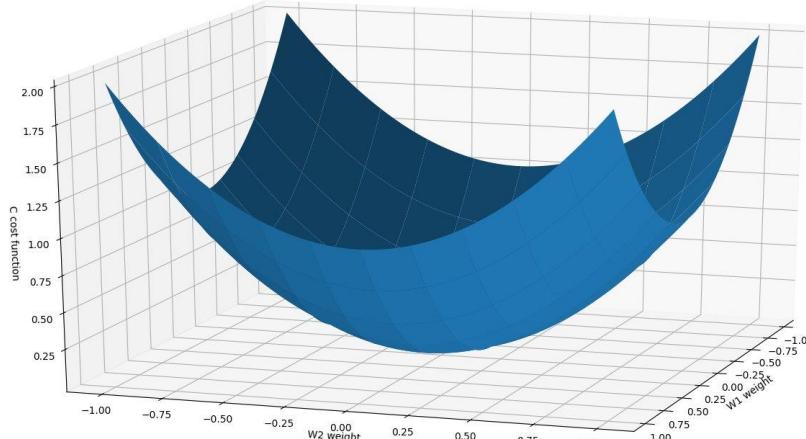
- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

- $y = f(x_1, x_2, \dots, x_k)$



# It's Calculus Time



- **Derivatives**


- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

- $y = f(x_1, x_2, \dots, x_k)$
- How a small change in  $x_i$  changes  $y$ , keeping all  $x_{k \neq i}$  constant



# It's Calculus Time



- **Derivatives**

- $y = f(x)$
- How a small change in  $x$  changes  $y$
- $\lim_{a \rightarrow 0} \frac{f(x+a)-f(x)}{a}$  or  $\frac{dy}{dx}$
- Slope of the tangent line

- **Gradient descent**

- Minimize a function by subtracting derivative at the point

- **Partial Derivative**

- $y = f(x_1, x_2, \dots, x_k)$
- How a small change in  $x_i$  changes  $y$ , keeping all  $x_{k \neq i}$  constant
- $\frac{\partial y}{\partial x}$  instead of  $\frac{dy}{dx}$

# The chain rule of derivatives

# The chain rule of derivatives

$$\frac{d}{dx} g(f(x)) = \frac{dg}{df} \frac{df}{dx}$$

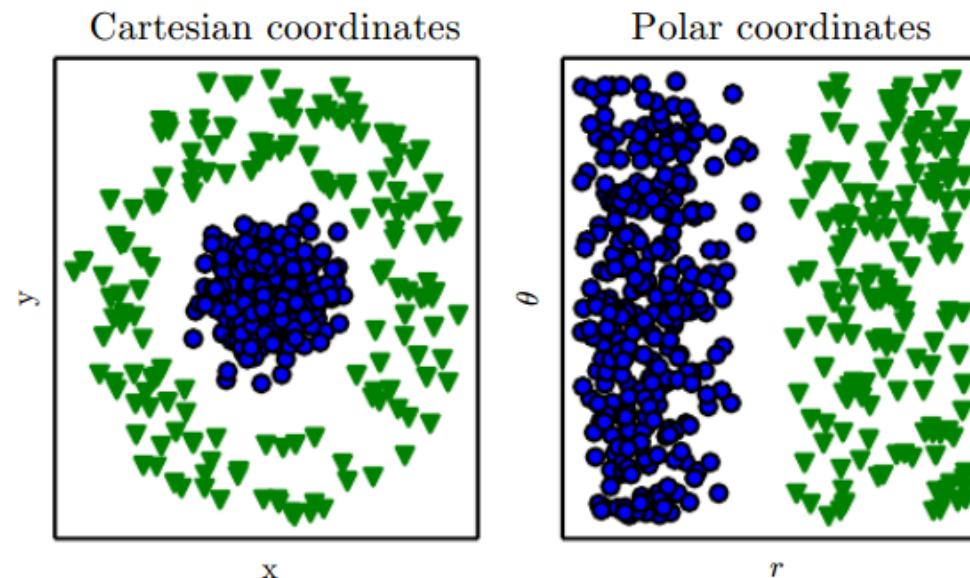
# Perceptron

# Perceptron

- Linear binary classifier

# Perceptron

- Linear binary classifier
  - Linearly separable data



# Perceptron

- Linear binary classifier
  - Linearly separable data
  - One out of two classes




# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning




# Perceptron

- Linear binary classifier
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- Components




# Perceptron

- Linear binary classifier
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- Components
  - Input values



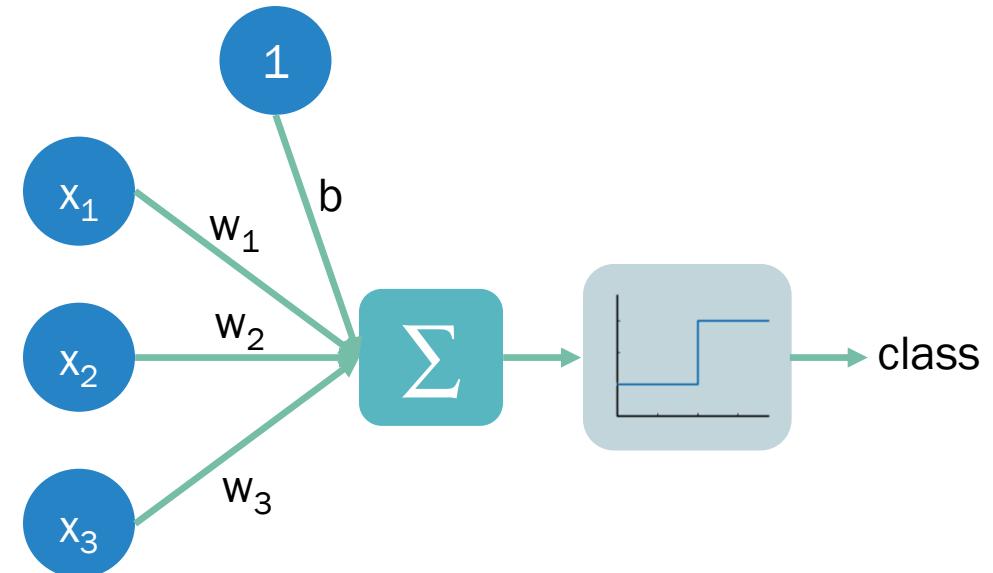
# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- **Components**
  - Input values
  - Weights and Biases




# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- **Components**
  - Input values
  - Weights and Biases
  - Sum




# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- **Components**
  - Input values
  - Weights and Biases
  - Sum
  - Activation function



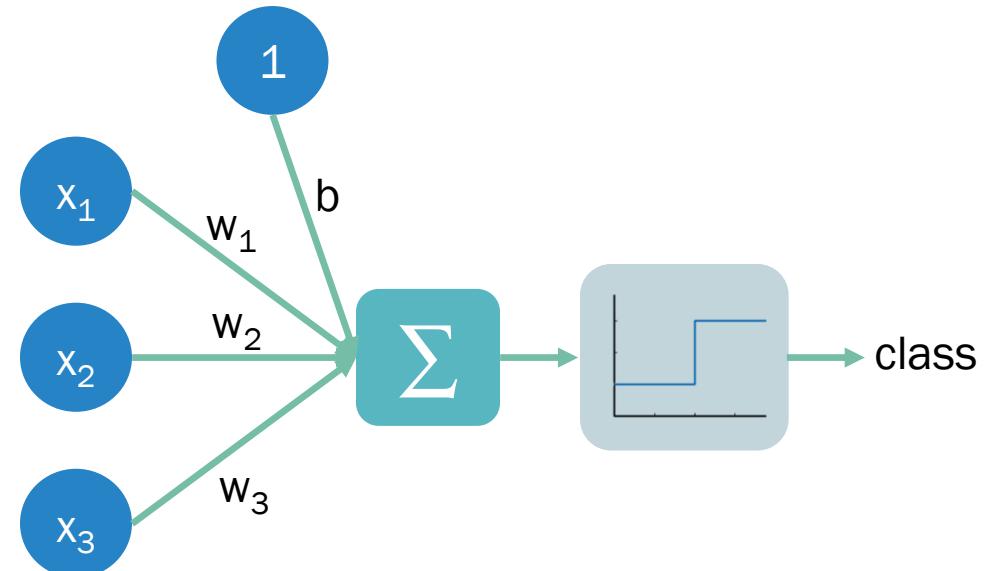
# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- **Components**
  - Input values
  - Weights and Biases
  - Sum
  - Activation function
  - Weight update from errors (*training*)



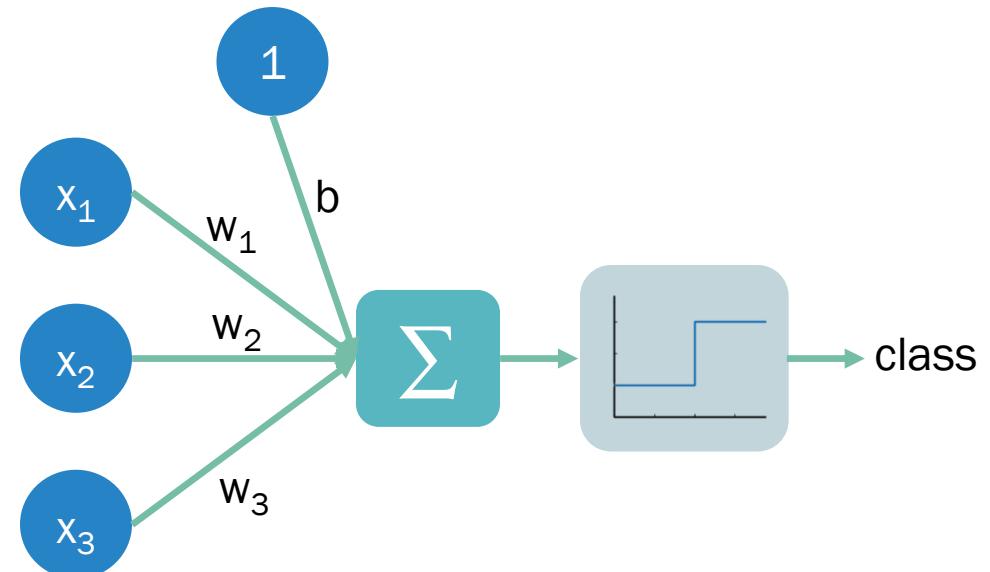
# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- **Components**
  - Input values
  - Weights and Biases
  - Sum
  - Activation function
  - Weight update from errors (*training*)
- **Steps**




# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- **Components**
  - Input values
  - Weights and Biases
  - Sum
  - Activation function
  - Weight update from errors (*training*)
- **Steps**
  1. Multiply inputs by the weights




# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- **Components**
  - Input values
  - Weights and Biases
  - Sum
  - Activation function
  - Weight update from errors (*training*)
- **Steps**
  1. Multiply inputs by the weights
  2. Add up the values




# Perceptron

- **Linear binary classifier**
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- **Components**
  - Input values
  - Weights and Biases
  - Sum
  - Activation function
  - Weight update from errors (*training*)
- **Steps**
  1. Multiply inputs by the weights
  2. Add up the values
  3. Apply the activation function



# Perceptron


- Linear binary classifier
  - Linearly separable data
  - One out of two classes
  - Supervised learning
- Components
  - Input values
  - Weights and Biases
  - Sum
  - Activation function
  - Weight update from errors (*training*)
- Steps
  1. Multiply inputs by the weights
  2. Add up the values
  3. Apply the activation function
  4. Update the weights given error (*training*)



# Step & Error functions

# Step & Error functions

- Sigmoid function




# Step & Error functions

- **Sigmoid function**

- $\frac{1}{1+e^{-x}}$

- Can be interpreted as probability



# Step & Error functions

- Sigmoid function
  - $\frac{1}{1+e^{-x}}$
  - Can be interpreted as probability
- Cross-entropy



# Step & Error functions

- **Sigmoid function**

- $\frac{1}{1+e^{-x}}$

- Can be interpreted as probability

- **Cross-entropy**

- $-\sum_i(t_i \ln(y_i) + (1 - t_i) \ln(1 - y_i))$
- Distance between two distributions



# Step & Error functions

- **Sigmoid function**

- $\frac{1}{1+e^{-x}}$

- Can be interpreted as probability

- **Cross-entropy**

- $-\sum_i(t_i \ln(y_i) + (1 - t_i) \ln(1 - y_i))$
- Distance between two distributions
- Rooted on information theory
  - Huffman coding



# The Internals

---

Is it all a black box?

---

*"For every action there is an equal and opposite reaction "*

*Newton's 3<sup>rd</sup> Law of Motion*


# Setup

Single datum


Binary classification  
problem

Objective:  
 $\arg \min_{w,b} \text{Error}$

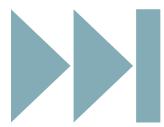
# Anatomy of a neural network



# Anatomy of a neural network



# Derivation Steps



Forward  
Propagation



Error  
Computation



Back  
Propagation



Parameters  
update

# Implementation

---

Does it even work?

---

*"Machines take me by surprise with great frequency "*

*Alan Turing*

# Just a tiny bit more of fundamentals

- Dot product of two matrices

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`?
```



# Just a tiny bit more of fundamentals



- Dot product of two matrices
  - Iterating over the arrays
  - 247 ms

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c_slow = np.zeros((100, 100))

for i in range(a.shape[0]):
    for j in range(b.shape[1]):
        acc = 0
        for k in range(a.shape[1]):
            acc += a[i, k] * b[k, j]
        c_slow[i, j] = acc
```

# Just a tiny bit more of fundamentals



- Dot product of two matrices
  - Iterating over the arrays
    - 247 ms
  - Numpy's dot product
    - 110 µs

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

# Just a tiny bit more of fundamentals



- Dot product of two matrices
  - Iterating over the arrays
    - 247 ms
  - Numpy's dot product
    - 110  $\mu$ s
- **Vectorization** – rewriting loops into efficient parallelizable operations

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

# Just a tiny bit more of fundamentals



- Dot product of two matrices
  - Iterating over the arrays
    - 247 ms
  - Numpy's dot product
    - 110 µs
- **Vectorization** – rewriting loops into efficient parallelizable operations
  - “Closer-to-the-metal” computations

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

# Just a tiny bit more of fundamentals



- Dot product of two matrices
  - Iterating over the arrays
    - 247 ms
  - Numpy's dot product
    - 110 µs
- **Vectorization** – rewriting loops into efficient parallelizable operations
  - “Closer-to-the-metal” computations
  - Prone to parallelization

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```

# Just a tiny bit more of fundamentals



- Dot product of two matrices
  - Iterating over the arrays
    - 247 ms
  - Numpy's dot product
    - 110 µs
- **Vectorization** – rewriting loops into efficient parallelizable operations
  - “Closer-to-the-metal” computations
  - Prone to parallelization
  - Cleaner code

```
import numpy as np

a = np.random.normal(size=(100, 50))
b = np.random.normal(size=(50, 100))

# dot product of `a` and `b`

c = a.dot(b)
```


**ENOUGH TALK**



**LET'S CODE**

memegenerator.net

<http://bit.ly/2GFMJJP>



# Today I learned

## What we did:

- Derived the math behind traditional neural networks for binary classification
- Implemented and evaluated them using Numpy

## Potential paths of exploration:

- Regression
  - Mean squared error
- Multiclass
  - Softmax activation
  - Cross entropy
- Refactor the implementation
- ...

# Bibliography

- [Deep Learning](#) by Goodfellow et al. for images and inspiration
- [Notes on Backpropagation](#)
- [Census income dataset](#) from [UCI Repository](#)
- [Visual Information Theory](#) by [Christopher Olah](#)
- [Neural Networks MOOC](#) by [Geoffrey Hinton](#)
- [What the Hell is a Perceptron](#) by [Sagar Sharma](#)

# Acknowledgements

- Icons made by [Freepik](#), [dDara](#), [surang](#), [Vectors Market](#), obtained from [Flaticon](#)

# Thank you



[@diogojapinto](https://twitter.com/diogojapinto)



[github.com/diogojapinto](https://github.com/diogojapinto)



[pt.linkedin.com/in/diogojapinto](https://pt.linkedin.com/in/diogojapinto)